



# Developer Training for Spark and Hadoop I

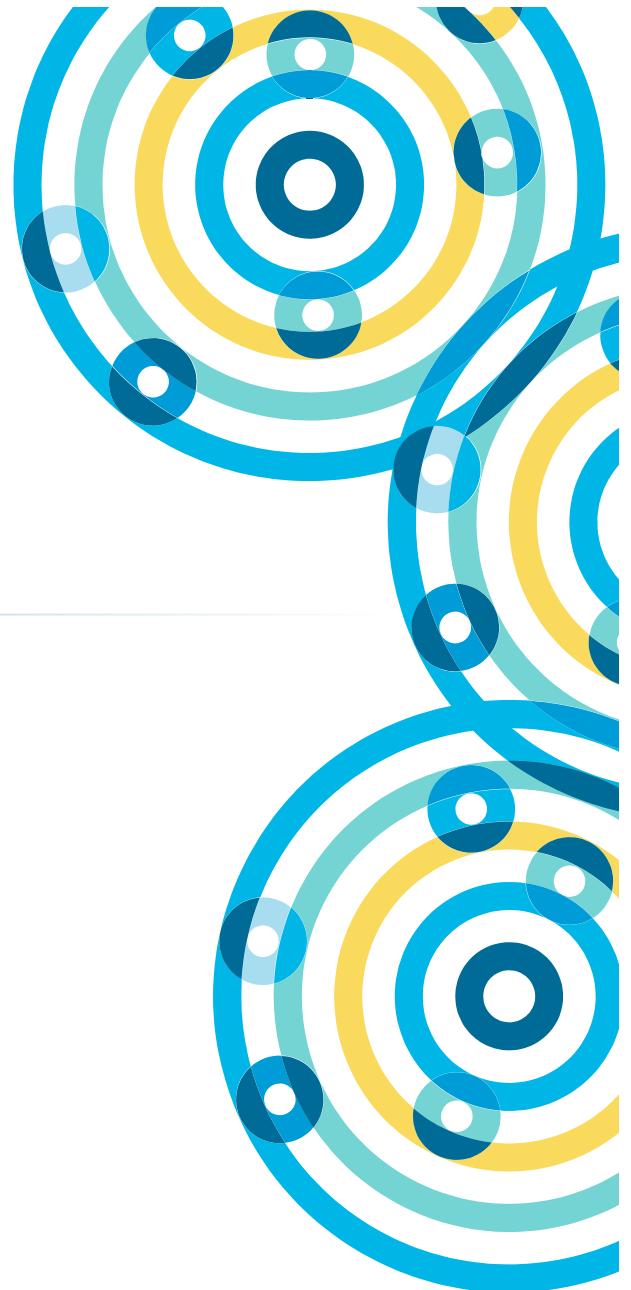




# Introduction to Impala and Hive

---

Chapter 5



# Course Chapters

- Introduction
- Introduction to Hadoop and the Hadoop Ecosystem
- Hadoop Architecture and HDFS
- Importing Relational Data with Apache Sqoop
- **Introduction to Impala and Hive**
- Working with Tables in Impala
- Data Formats
- Data File Partitioning
- Capturing Data with Apache Flume
- Spark Basics
- Working with RDDs in Spark
- Aggregating Data with Pair RDDs
- Writing and Deploying Spark Applications
- Parallel Processing in Spark
- Spark RDD Persistence
- Common Patterns in Spark Data Processing
- Spark SQL and DataFrames
- Conclusion

Course Introduction

Introduction to Hadoop

## Importing and Modeling Structured Data

Ingesting Streaming Data

Distributed Data Processing with Spark

Course Conclusion

# Introduction to Impala and Hive

---

**In this chapter you will learn**

- **What Hive is**
- **What Impala is**
- **How Impala and Hive Compare**
- **How to query data using Impala and Hive**
- **How Hive and Impala differ from a relational database**
- **Ways in which organizations use Hive and Impala**

# Chapter Topics

## Introduction to Impala and Hive

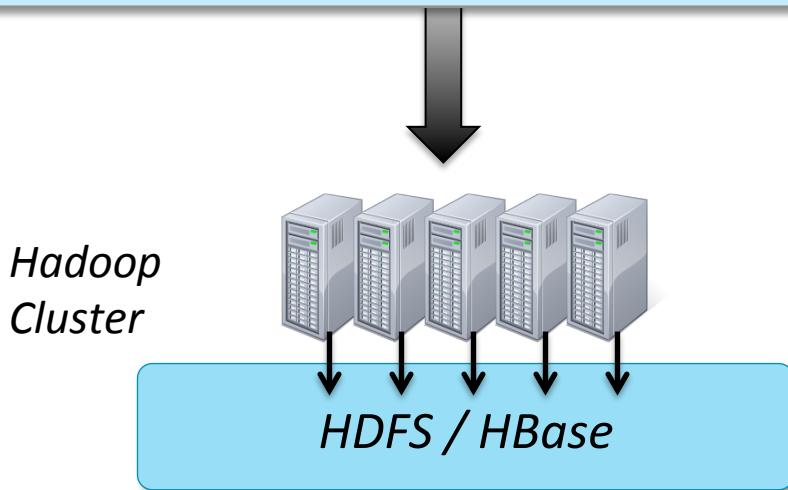
## Importing and Modeling Structured Data

- **Introduction to Impala and Hive**
- Why Use Impala and Hive?
- Querying Data With Impala and Hive
- Comparing Hive and Impala to Traditional Databases

## Introduction to Impala and Hive (1)

- Impala and Hive are both tools that provide SQL querying of data stored in HDFS / HBase

```
SELECT zipcode, SUM(cost) AS total  
FROM customers  
JOIN orders  
ON (customers.cust_id = orders.cust_id)  
WHERE zipcode LIKE '63%'  
GROUP BY zipcode  
ORDER BY total DESC;
```



## Introduction to Impala and Hive (2)

---

- **Apache Hive is a high-level abstraction on top of MapReduce**
  - Uses HiveQL
  - Generates MapReduce or Spark\* jobs that run on the Hadoop cluster
  - Originally developed at Facebook around 2007
    - Now an open-source Apache project
- **Cloudera Impala is a high-performance dedicated SQL engine**
  - Uses Impala SQL
  - Inspired by Google's Dremel project
  - Query latency measured in milliseconds
  - Developed at Cloudera in 2012
    - Open-source with an Apache license



\* Hive-on-Spark is currently in beta testing

# What's the Difference?

---

- **Hive has more features**

- E.g. Complex data types (arrays, maps) and full support for windowing analytics
  - Highly extensible
  - Commonly used for batch processing

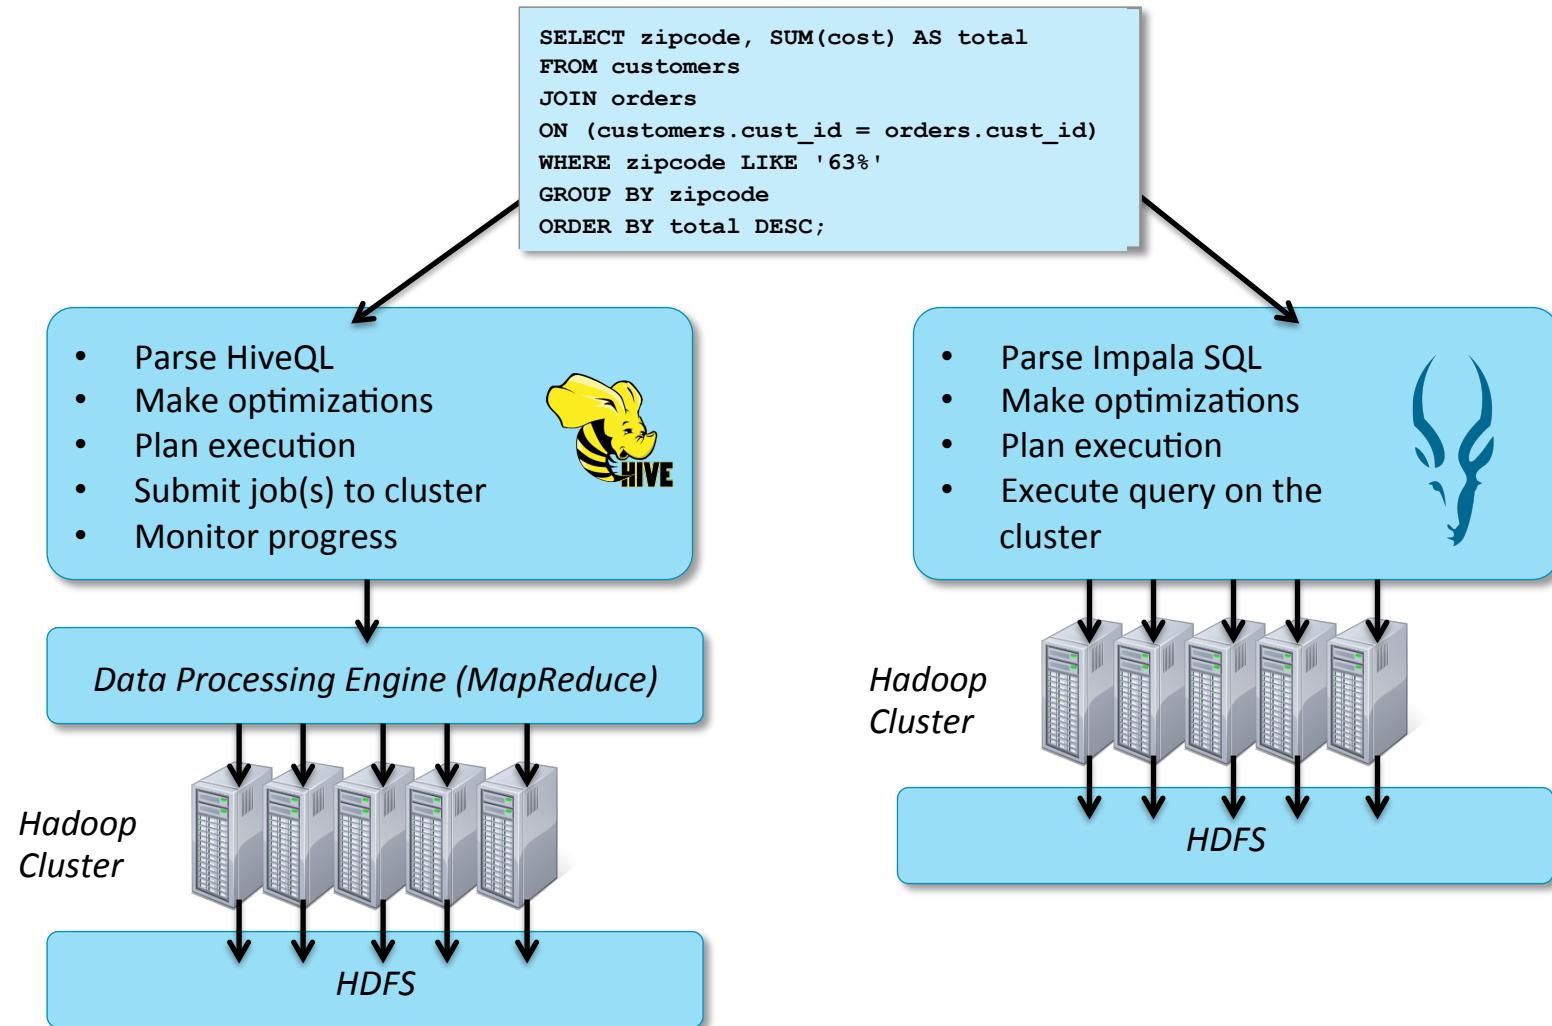


- **Impala is much faster**

- Specialized SQL engine offers 5x to 50x better performance
  - Ideal for interactive queries and data analysis
  - More features being added over time



# High-Level Overview



# Chapter Topics

## Introduction to Impala and Hive

## Importing and Modeling Structured Data

- Introduction to Impala and Hive
- **Why Use Impala and Hive?**
- Querying Data With Impala and Hive
- Comparing Hive to Traditional Databases
- Conclusion

## Why Use Hive and Impala?

---

- **Brings large-scale data analysis to a broader audience**
  - No software development experience required
  - Leverage existing knowledge of SQL
- **More productive than writing MapReduce or Spark directly**
  - Five lines of HiveQL/Impala SQL might be equivalent to 200 lines or more of Java
- **Offers interoperability with other systems**
  - Extensible through Java and external scripts
  - Many business intelligence (BI) tools support Hive and/or Impala

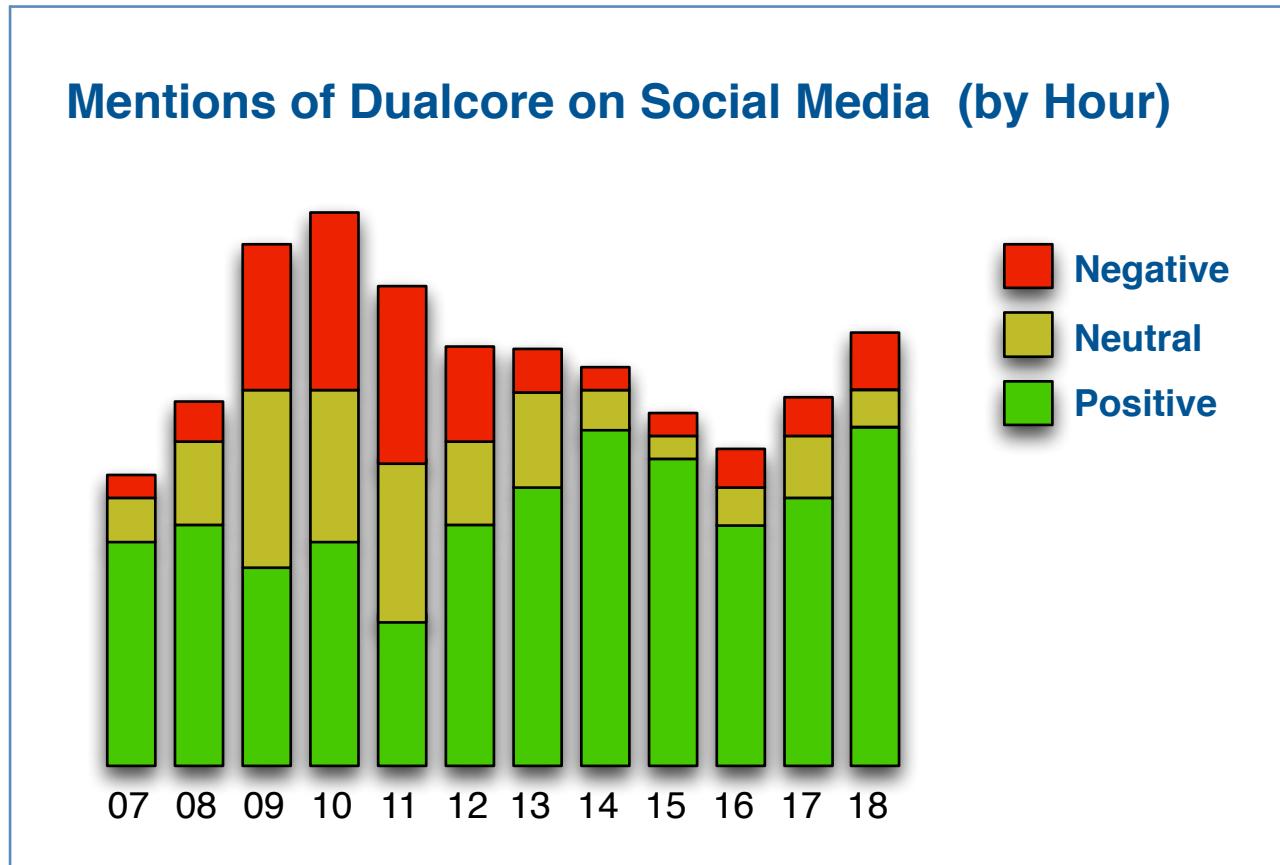
## Use Case: Log File Analytics

- Server log files are an important source of data
- Hive and Impala allow you to treat a directory of log files like a table
  - Allows SQL-like queries against raw data

Dualcore Inc. Public Web Site (June 1 - 8)					
Product	Unique Visitors	Page Views	Average Time on Page	Bounce Rate	Conversion Rate
Tablet	5,278	5,894	17 seconds	23%	65%
Notebook	4,139	4,375	23 seconds	47%	31%
Stereo	2,873	2,981	42 seconds	61%	12%
Monitor	1,749	1,862	26 seconds	74%	19%
Router	987	1,139	37 seconds	56%	17%
Server	314	504	53 seconds	48%	28%
Printer	86	97	34 seconds	27%	64%

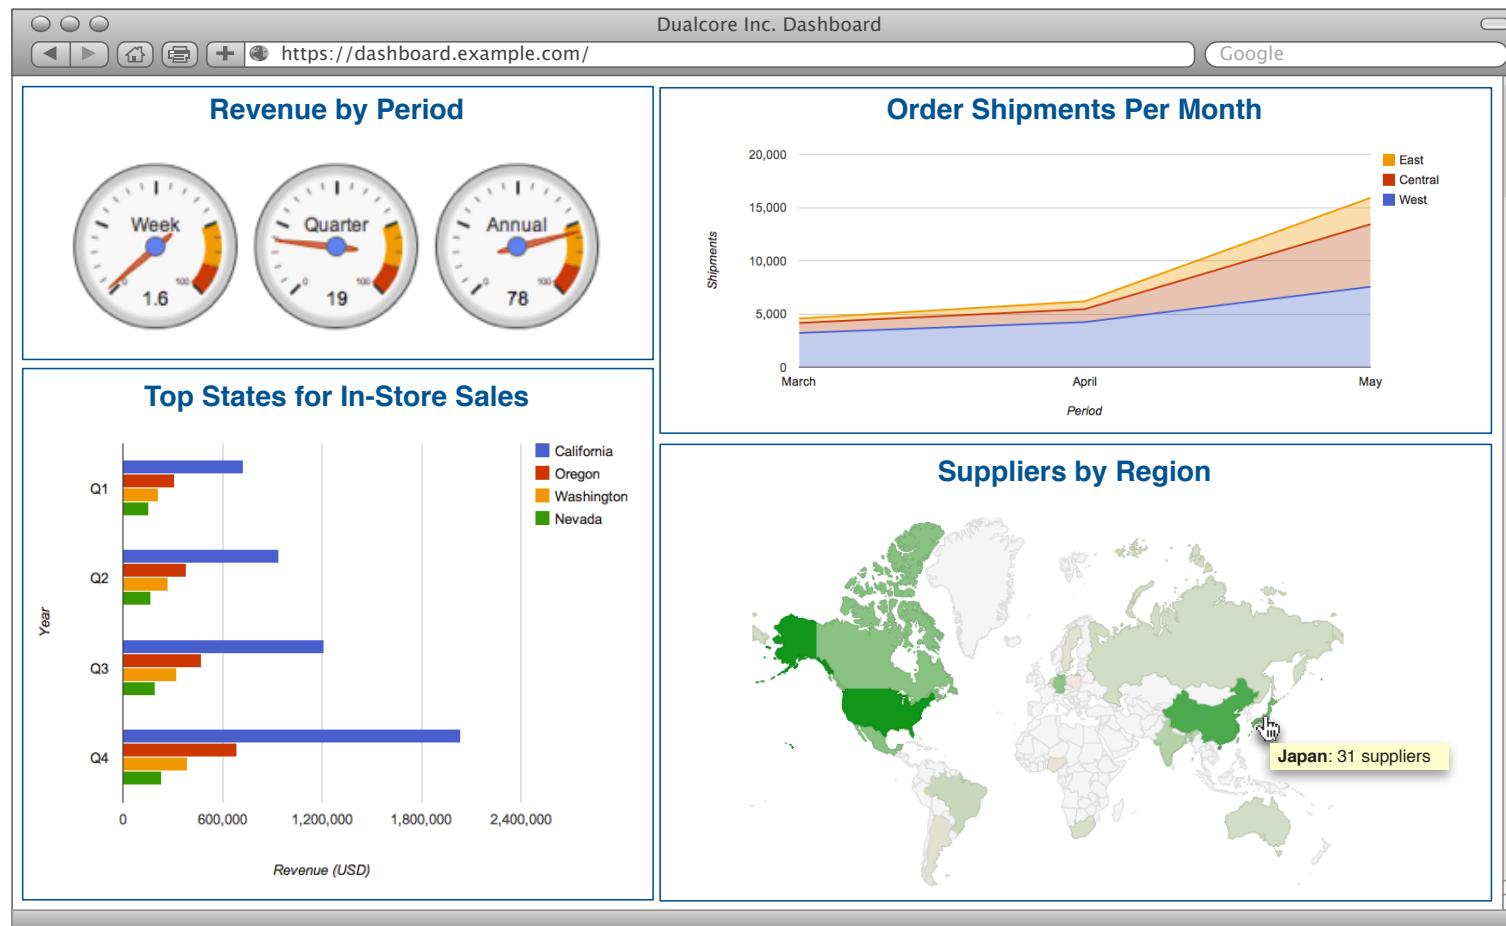
## Use Case: Sentiment Analysis

- Many organizations use Hive or Impala to analyze social media coverage



# Use Case: Business Intelligence

- Many leading business intelligence tools support Hive and Impala



# Chapter Topics

## Introduction to Impala and Hive

## Importing and Modeling Structured Data

- Introduction to Impala and Hive
- Why Use Impala and Hive?
- **Querying Data With Hive and Impala**
- Comparing Hive to Traditional Databases
- Conclusion

## Interacting with Hive and Impala

---

- **Hive and Impala offer many interfaces for running queries**
  - Command-line shell
    - Impala: Impala shell
    - Hive: Beeline
  - Hue Web UI
    - Hive Query Editor
    - Impala Query Editor
    - Metastore Manager
  - ODBC / JDBC

## Starting the Impala Shell

- You can execute statements in the Impala shell
  - This interactive tool is similar to the shell in MySQL
- Execute the `impala-shell` command to start the shell
  - Some log messages truncated to better fit the slide

```
$ impala-shell
Connected to localhost.localdomain:21000
Server version: impalad version 2.1.0-cdh5 (...)

Welcome to the Impala shell.
[localhost.localdomain:21000] >
```

- Use `-i hostname:port` option to connect to a different server

```
$ impala-shell -i myserver.example.com:21000
[myserver.example.com:21000] >
```

## Using the Impala Shell

---

- **Enter semicolon-terminated statements at the prompt**
  - Hit [Enter] to execute a query or command
  - Use the `quit` command to exit the shell
- **Use `impala-shell --help` for a full list of options**

## Executing Queries in the Impala Shell

```
> SELECT lname, fname FROM customers WHERE state = 'CA'  
limit 50;  
  
Query: select lname, fname FROM customers WHERE state =  
'CA' limit 50  
+-----+-----+  
| lname      | fname       |  
+-----+-----+  
| Ham        | Marilyn    |  
| Franks     | Gerard     |  
| Preston    | Mason      |  
| Cortez     | Pamela     |  
| ...         |            |  
| Falgoust   | Jennifer   |  
+-----+-----+  
Returned 50 row(s) in 0.17s  
  
>
```

Note: shell prompt abbreviated as >

## Interacting with the Operating System

- Use **shell** to execute system commands from within Impala shell

```
> shell date;  
Mon May 20 16:44:35 PDT 2013
```

- No direct support for HDFS commands

- But could run hdfs dfs using shell

```
> shell hdfs dfs -mkdir /reports/sales/2013;
```

## Running Impala Queries from the Command Line

- You can execute a file containing queries using the **-f** option

```
$ impala-shell -f myquery.sql
```

- Run queries directly from the command line with the **-q** option

```
$ impala-shell -q 'SELECT * FROM users'
```

- Use **-o** to capture output to file

- Optionally specify delimiter

```
$ impala-shell -f myquery.sql \
-o results.txt \
--delimited \
--output_delimiter=','
```

## Starting Beeline (Hive's Shell)

- You can execute HiveQL statements in the Beeline shell
  - Interactive shell based on the SQLLine utility
  - Similar to the Impala shell
- Start Beeline by specifying the URL for a Hive2 server
  - Plus username and password if required

```
$ beeline -u jdbc:hive2://host:10000 \
-n username -p password

0: jdbc:hive2://localhost:10000>
```

## Executing Queries in Beeline

- SQL commands are terminated with semi-colon ( ; )
- Similar to Impala shell
  - Results formatting is slightly different

```
1: url> SELECT lname, fname FROM customers
. . . > WHERE state = 'CA' LIMIT 50;

+-----+-----+
| lname | fname |
+-----+-----+
| Ham   | Marilyn |
| Franks | Gerard |
| Preston | Mason |
...
| Falgoust | Jennifer |
+-----+-----+
50 rows selected (15.829 seconds)

1: url>
```

## Using Beeline

---

- Execute Beeline commands with ‘!’
  - No terminator character
- Some commands
  - **!connect url** – connect to a different Hive2 server
  - **!exit** – exit the shell
  - **!help** – show the full list of commands
  - **!verbose** – show added details of queries

```
0: jdbc:hive2://localhost:10000> !exit
```

## Executing Hive Queries from the Command Line

- You can also execute a file containing HiveQL code using the **-f** option

```
$ beeline -u ... -f myquery.hql
```

- Or use HiveQL directly from the command line using the **-e** option

```
$ beeline -u ... -e 'SELECT * FROM users'
```

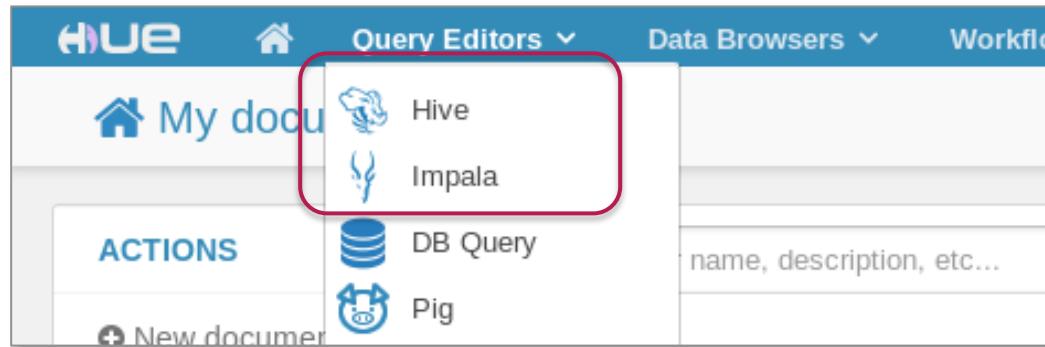
- Use the **--silent** option to suppress informational messages
  - Can also be used with **-e** or **-f** options

```
$ beeline -u ... --silent
```

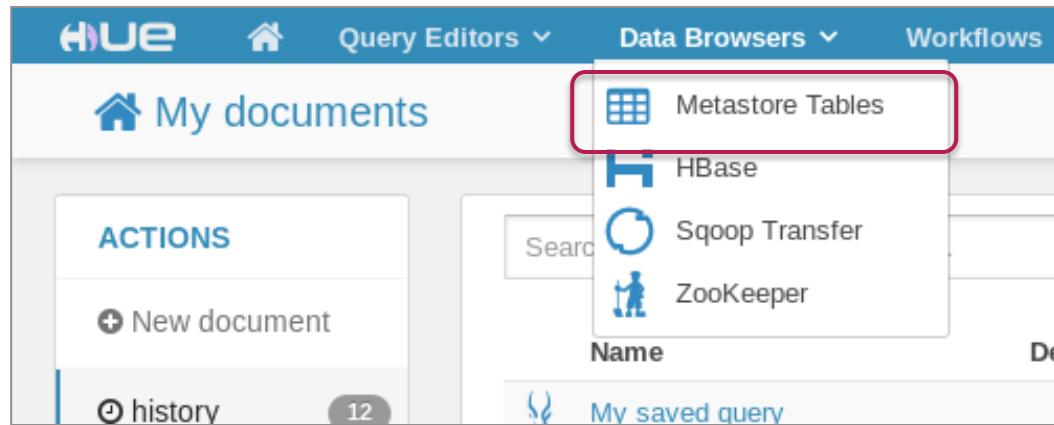
# Using Hue with Hive and Impala

You can use Hue to...

Query data with  
Hive or Impala



View and manage  
the Metastore



# The Hue Query Editor

- The Impala and Hive Query editors are nearly identical

The screenshot shows the Hue Query Editor interface. On the left, there's a sidebar for database selection and schema exploration. A blue box labeled "Choose a database" points to the "DATABASE..." dropdown set to "default". Another blue box labeled "Explore schema and sample data" points to the "customers" table schema, which includes columns: cust\_id (int), fname (str), lname (str), address (s), city (string), state (str), and zipcode (s). The main area is the "Query Editor" tab, where a query is entered in the text area:

```
1 SELECT * FROM customers WHERE state = 'CA';
```

A blue callout box with the text "Enter, edit, save and execute queries" points to this area. Below the query text are buttons for "Execute", "Save as...", "Explain", and "or create a New query". The results section at the bottom shows a table of data with columns: cust\_id, fname, lname, address, city, state, and zipcode. The data rows are:

	cust_id	fname	lname	address	city	state	zipcode
0	1000002	Marilyn	Ham	25831 North 25th Street	Concord	CA	94522
1	1000006	Gerard	Franks	356 Turner Street	Pioneer	CA	95666
2	1000010	Mason	Preston	2656 West 13th Street	Redwood Valley	CA	95470
3	1000012	Pamela	Cortez	2279 North Mulberry Avenue	San Francis		

A blue callout box with the text "View results, logs, reports, etc." points to the "Results" tab in the navigation bar.

# Chapter Topics

## Introduction to Impala and Hive

## Importing and Modeling Structured Data

- Introduction to Impala and Hive
- Why Use Impala and Hive?
- Querying Data With Impala and Hive
- **Comparing Hive and Impala to Traditional Databases**
- Conclusion

## Your Cluster is Not a Database Server

---

- **Client-server database management systems have many strengths**
  - Very fast response time
  - Support for transactions
  - Allow modification of existing records
  - Can serve thousands of simultaneous clients
- **Your Hadoop cluster is not an RDBMS**
  - Hive generates processing engine jobs (MapReduce) from HiveQL queries
    - Limitations of HDFS and MapReduce still apply
  - Impala is faster but not intended for the throughput speed required for an OLTP database
  - No transaction support

## Comparing Hive and Impala To A Relational Database

---

	Relational Database	Hive	Impala
<b>Query language</b>	SQL (full)	SQL (subset)	SQL (subset)
<b>Update individual records</b>	Yes	No	No
<b>Delete individual records</b>	Yes	No	No
<b>Transactions</b>	Yes	No	No
<b>Index support</b>	Extensive	Limited	No
<b>Latency</b>	Very low	High	Low
<b>Data size</b>	Terabytes	Petabytes	Petabytes

# Chapter Topics

## Introduction to Impala and Hive

## Importing and Modeling Structured Data

- Introduction to Impala and Hive
- Why Use Impala and Hive?
- Querying Data With Impala and Hive
- Comparing Hive and Impala to Traditional Databases
- **Conclusion**

## Essential Points

---

- **Impala and Hive are tools for performing SQL queries on data in HDFS**
- **HiveQL and Impala SQL are very similar to SQL-92**
  - Easy to learn for those with relational database experience
  - However, does *not* replace your RDBMS
- **Hive generates jobs that run on the Hadoop cluster data processing engine**
  - Runs MapReduce jobs on Hadoop based on HiveQL statements
- **Impala execute queries directly on the Hadoop cluster**
  - Uses a very fast specialized SQL engine, not MapReduce

## Bibliography

---

The following offer more information on topics discussed in this chapter

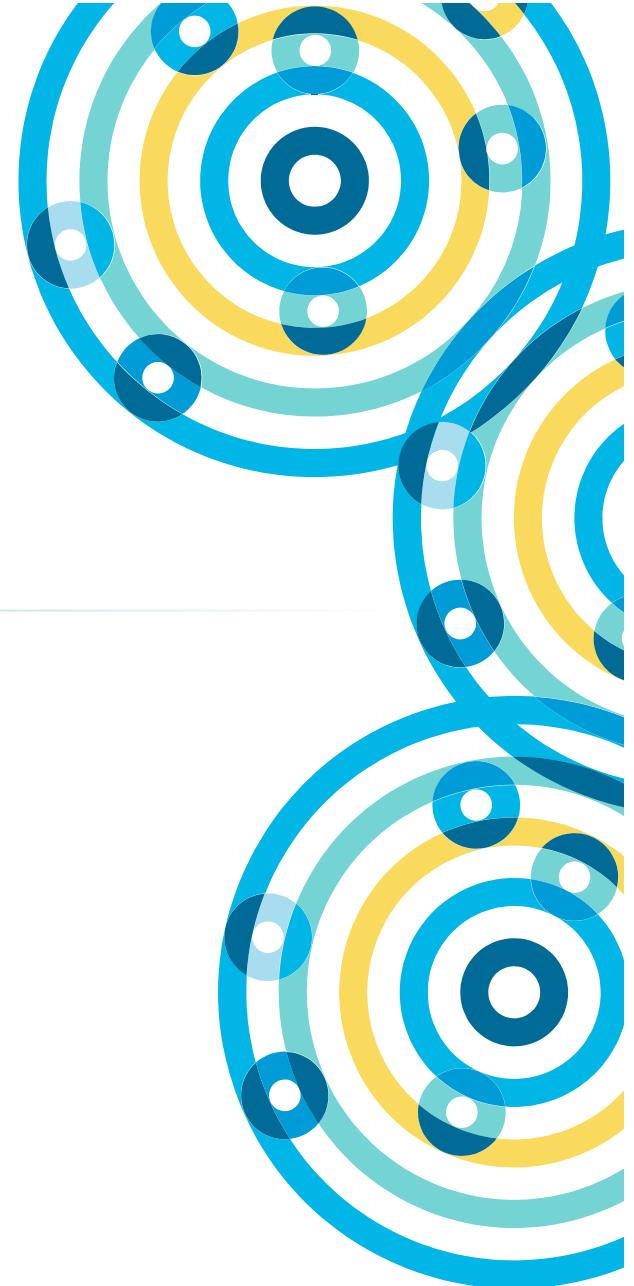
- ***Cloudera Impala (free O'Reilly book)***
  - <http://tiny.cloudera.com/impalabook>
- ***Programming Hive (O'Reilly book)***
  - <http://tiny.cloudera.com/programminghive>
- **Data Analysis with Hadoop and Hive at Orbitz**
  - <http://tiny.cloudera.com/dac09b>
- **Sentiment Analysis Using Apache Hive**
  - <http://tiny.cloudera.com/dac09c>
- ***Wired Article on Impala***
  - <http://tiny.cloudera.com/wiredimpala>



# Modeling and Managing Data with Impala and Hive

---

Chapter 6



# Course Chapters

- Introduction
- Introduction to Hadoop and the Hadoop Ecosystem
- Hadoop Architecture and HDFS
- Importing Relational Data with Apache Sqoop
- Introduction to Impala and Hive
- **Modeling and Managing Data with Impala and Hive**
- Data Formats
- Data File Partitioning
- Capturing Data with Apache Flume
- Spark Basics
- Working with RDDs in Spark
- Aggregating Data with Pair RDDs
- Writing and Deploying Spark Applications
- Parallel Processing in Spark
- Spark RDD Persistence
- Common Patterns in Spark Data Processing
- Spark SQL and DataFrames
- Conclusion

Course Introduction

Introduction to Hadoop

## Importing and Modeling Structured Data

Ingesting Streaming Data

Distributed Data Processing with Spark

Course Conclusion

# Modeling and Managing Data in Impala and Hive

---

**In this chapter you will learn**

- **How Impala and Hive use the Metastore**
- **How to use Impala SQL and HiveQL DDL to create tables**
- **How to create and manage tables using Hue or HCatalog**
- **How to load data into tables using Impala, Hive, or Sqoop**

# Chapter Topics

**Modeling and Managing Data With Impala and Hive**

**Importing and Modeling Structured Data**

- **Data Storage Overview**

- Creating Databases and Tables
- Loading Data into Tables
- HCatalog
- Impala Metadata Caching
- Conclusion
- Hands-On Exercise: Create and Populate Tables in Impala or Hive

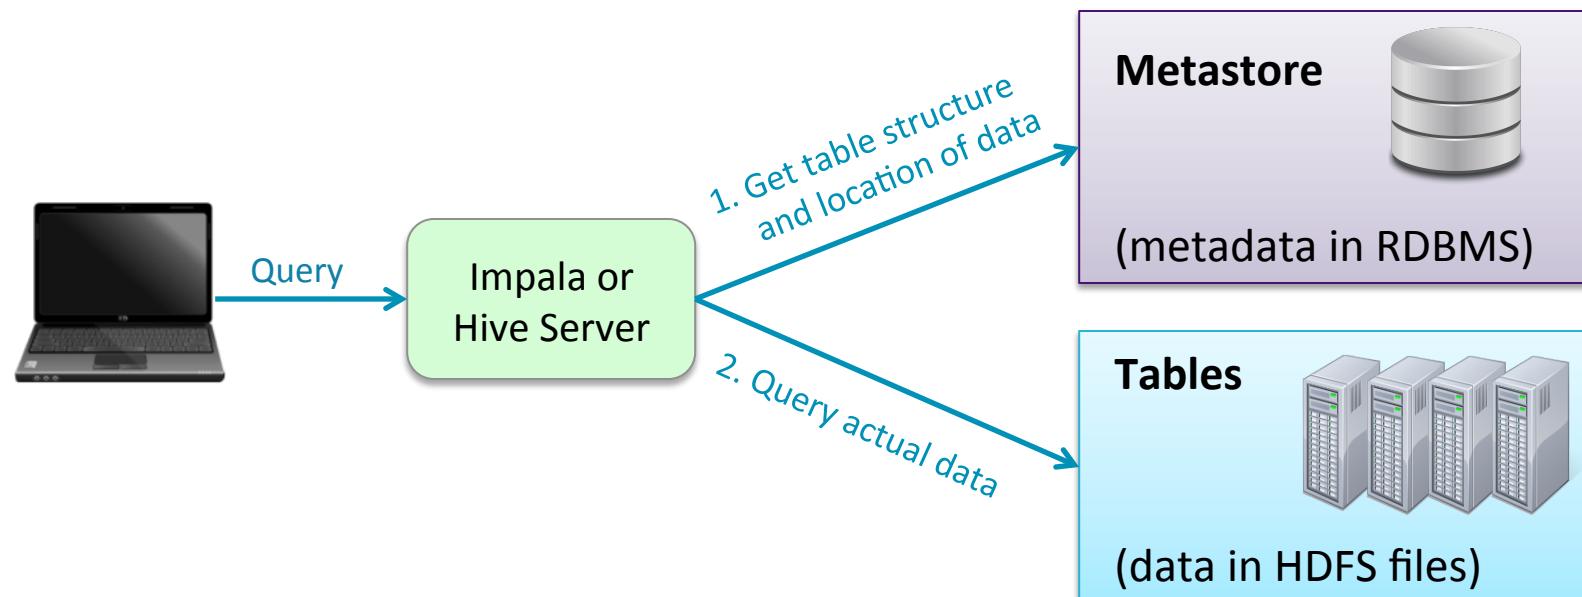
## How Hive and Impala Load and Store Data (1)

---

- **Queries operate on tables, just like in an RDBMS**
  - A table is simply an HDFS directory containing one or more files
  - Default path: `/user/hive/warehouse/<table_name>`
  - Supports many formats for data storage and retrieval
- **What is the structure and location of tables?**
  - These are specified when tables are created
  - This metadata is stored in the *Metastore*
    - Contained in an RDBMS such as MySQL
- **Hive and Impala work with the same data**
  - Tables in HDFS, metadata in the Metastore

## How Hive and Impala Load and Store Data (2)

- **Hive and Impala use the Metastore to determine data format and location**
  - The query itself operates on data stored in HDFS



## Data and Metadata

- **Data refers to the information you store and process**
  - Billing records, sensor readings, and server logs are examples of data
- **Metadata describes important aspects of that data**
  - Field name and order are examples of metadata

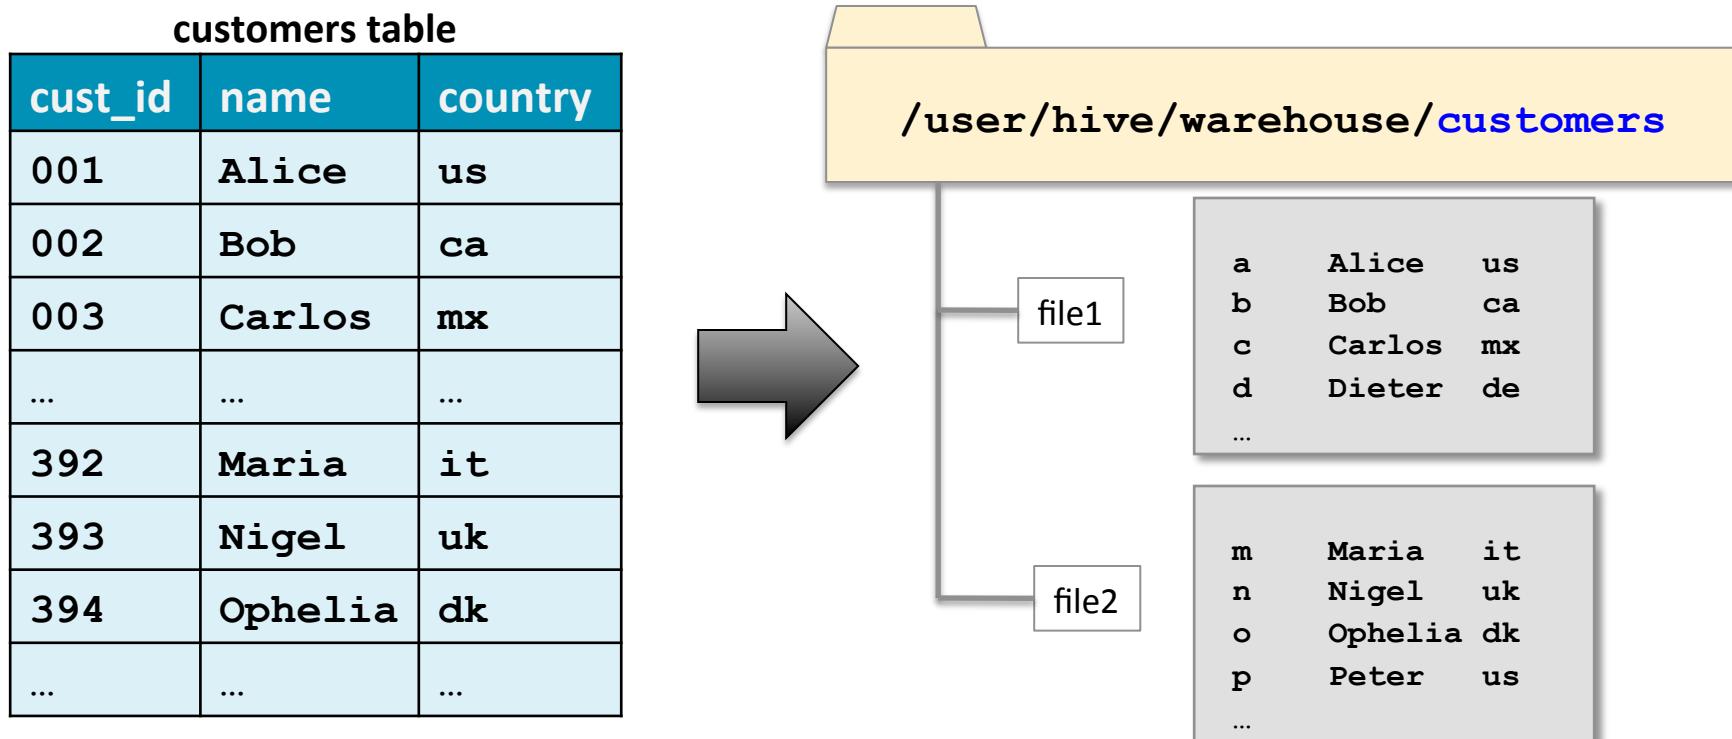


The diagram illustrates the relationship between metadata and data. A vertical bracket on the left side groups the labels "Metadata" and "Data" with the table to their right. The table consists of a header row and eight data rows, each containing three columns: "cust\_id", "name", and "country".

cust_id	name	country
001	Alice	us
002	Bob	ca
003	Carlos	mx
...	...	...
392	Maria	it
393	Nigel	uk
394	Ophelia	dk
...	...	...

# The Data Warehouse Directory

- By default, data is stored in the HDFS directory `/user/hive/warehouse`
- Each table is a subdirectory containing any number of files



# Chapter Topics

**Modeling and Managing Data With Impala and Hive**

**Importing and Modeling Structured Data**

- Data Storage Overview
- **Creating Databases and Tables**
- Loading Data into Tables
- HCatalog
- Impala Metadata Caching
- Conclusion
- Hands-On Exercise: Create and Populate Tables in Impala or Hive

## Defining Databases and Tables

---

- **Databases and tables are created and managed using the DDL (Data Definition Language) of HiveQL or Impala SQL**
  - Very similar to standard SQL DDL
  - Some minor differences between Hive and Impala DDL will be noted

## Creating a Database

---

- **Hive and Impala databases are simply namespaces**
  - Helps to organize your tables
- **To create a new database**

```
CREATE DATABASE loudacre;
```

1. Adds the database definition to the metastore
2. Creates a storage directory in HDFS  
e.g./user/hive/warehouse/loudacre.db

- **To conditionally create a new database**
  - Avoids error in case database already exists (useful for scripting)

```
CREATE DATABASE IF NOT EXISTS loudacre;
```

## Removing a Database

- Removing a database is similar to creating it
  - Just replace **CREATE** with **DROP**

```
DROP DATABASE loudacre;
```

```
DROP DATABASE IF EXISTS loudacre;
```

- These commands will fail if the database contains tables
  - In Hive: Add the **CASCADE** keyword to force removal
  - Caution: this command might remove data in HDFS!



```
DROP DATABASE loudacre CASCADE;
```

# Data Types

---

- **Each column is assigned a specific data type**
  - These are specified when the table is created
  - NULL values are returned for non-conforming data in HDFS
- **Here are some common data types**

Name	Description	Example Value
<b>STRING</b>	Character data (of any length)	<b>Alice</b>
<b>BOOLEAN</b>	<b>TRUE</b> or <b>FALSE</b>	<b>TRUE</b>
<b>TIMESTAMP</b>	Instant in time	<b>2014-03-14 17:01:29</b>
<b>INT</b>	Range: same as Java <b>int</b>	<b>84127213</b>
<b>BIGINT</b>	Range: same as Java <b>long</b>	<b>7613292936514215317</b>
<b>FLOAT</b>	Range: same as Java <b>float</b>	<b>3.14159</b>
<b>DOUBLE</b>	Range: same as Java <b>double</b>	<b>3.1415926535897932385</b>



Hive (not Impala) also supports a few complex types such as maps and arrays

## Creating a Table (1)

---

- Basic syntax for creating a table:

```
CREATE TABLE tablename (colname DATATYPE, ...)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY char  
STORED AS {TEXTFILE|SEQUENCEFILE|...}
```

- Creates a subdirectory in the database's **warehouse** directory in HDFS
  - Default database:  
**/user/hive/warehouse/*tablename***
  - Named database:  
**/user/hive/warehouse/*dbname*.db/*tablename***

## Creating a Table (2)

---

```
CREATE TABLE tablename (colname DATATYPE, ...)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS
```

```
STORED
```

Specify a name for the table, and list the column names and datatypes (see later)

## Creating a Table (3)

```
CREATE TABLE tablename (colname DATATYPE, ...)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY char
```

```
STORED AS TEXTFILE
```

This line states that fields in each file in the table's directory are delimited by some character. The default delimiter is Control-A, but you may specify an alternate delimiter...

## Creating a Table (4)

---

```
CREATE TABLE tablename (colname DATATYPE, ...)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY char  
STORED AS {TEXTFILE|SEQUENCEFILE|...}
```

...for example, tab-delimited data would require that you specify **FIELDS TERMINATED BY '\t'**

## Creating a Table (5)

---

```
CREATE TABLE tablename (colname DATATYPE, ...)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY char  
STORED AS {TEXTFILE|SEQUENCEFILE|...}
```

Finally, you may declare the file format. **STORED AS TEXTFILE** is the default and does not need to be specified.  
Other formats will be discussed later in the course.

## Example Table Definition

- The following example creates a new table named **jobs**
  - Data stored as text with four comma-separated fields per line

```
CREATE TABLE jobs (
    id INT,
    title STRING,
    salary INT,
    posted TIMESTAMP
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';
```

- Example of corresponding record for the table above

```
1,Data Analyst,100000,2013-06-21 15:52:03
```

## Creating Tables Based On Existing Schema

---

- Use **LIKE** to create a new table based on an existing table definition

```
CREATE TABLE jobs_archived LIKE jobs;
```

- Column definitions and names are derived from the existing table
  - New table will contain no data

## Creating Tables Based On Existing Data

- **Create a table based on a SELECT statement**
  - Often know as ‘Create Table As Select’ (CTAS)

```
CREATE TABLE ny_customers AS
    SELECT cust_id, fname, lname
        FROM customers
    WHERE state = 'NY';
```

- **Column definitions are derived from the existing table**
- **Column names are inherited from the existing names**
  - Use aliases in the SELECT statement to specify new names
- **New table will contain the selected data**

## Controlling Table Data Location

- By default, table data is stored in the warehouse directory
- This is not always ideal
  - Data might be shared by several users
- Use LOCATION to specify the directory where table data resides

```
CREATE TABLE jobs (
    id INT,
    title STRING,
    salary INT,
    posted TIMESTAMP
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/loudacre/jobs';
```

## Externally Managed Tables

- **CAUTION: Dropping a table removes its data in HDFS**
  - Tables are “managed” or “internal” by default
- **Using EXTERNAL when creating the table avoids this behavior**
  - Dropping an *external* table removes only its *metadata*

```
CREATE EXTERNAL TABLE adclicks
( campaign_id STRING,
  click_time TIMESTAMP,
  keyword STRING,
  site STRING,
  placement STRING,
  was_clicked BOOLEAN,
  cost SMALLINT)
LOCATION '/loudacre/ad_data' ;
```

## Exploring Tables (1)

- The **SHOW TABLES** command lists all tables in the current database

```
SHOW TABLES;
+-----+
| tab_name |
+-----+
| accounts |
| employees |
| job       |
| vendors   |
+-----+
```

- The **DESCRIBE** command lists the fields in the specified table

```
DESCRIBE jobs;
+-----+-----+-----+
| name | type  | comment |
+-----+-----+-----+
| id   | int   |          |
| title | string |          |
| salary | int   |          |
| posted | timestamp |          |
+-----+-----+-----+
```

## Exploring Tables (2)

- **DESCRIBE FORMATTED** also shows table properties

```
DESCRIBE FORMATTED jobs;
+-----+-----+-----+
| name      | type       | comment |
+-----+-----+-----+
| # col_name | data_type   | comment |
| id         | int         | NULL    |
| title      | string      | NULL    |
| salary     | int         | NULL    |
| posted     | timestamp   | NULL    |
|             | NULL        | NULL    |
| # Detailed Table | NULL      | NULL    |
| Information          |           |
| Database:            | default    | NULL    |
| Owner:              | training   | NULL    |
| CreateTime:          | Wed Jun 17 09:41:23 PDT 2015 | NULL    |
| LastAccessTime:      | UNKNOWN    | NULL    |
| Protect Mode:        | None       | NULL    |
| Retention:           | 0          | NULL    |
| Location:            | hdfs://localhost:8020/loudacre/jobs | NULL    |
| Table Type:          | MANAGED_TABLE | NULL    |
...
...
```

## Exploring Tables (3)

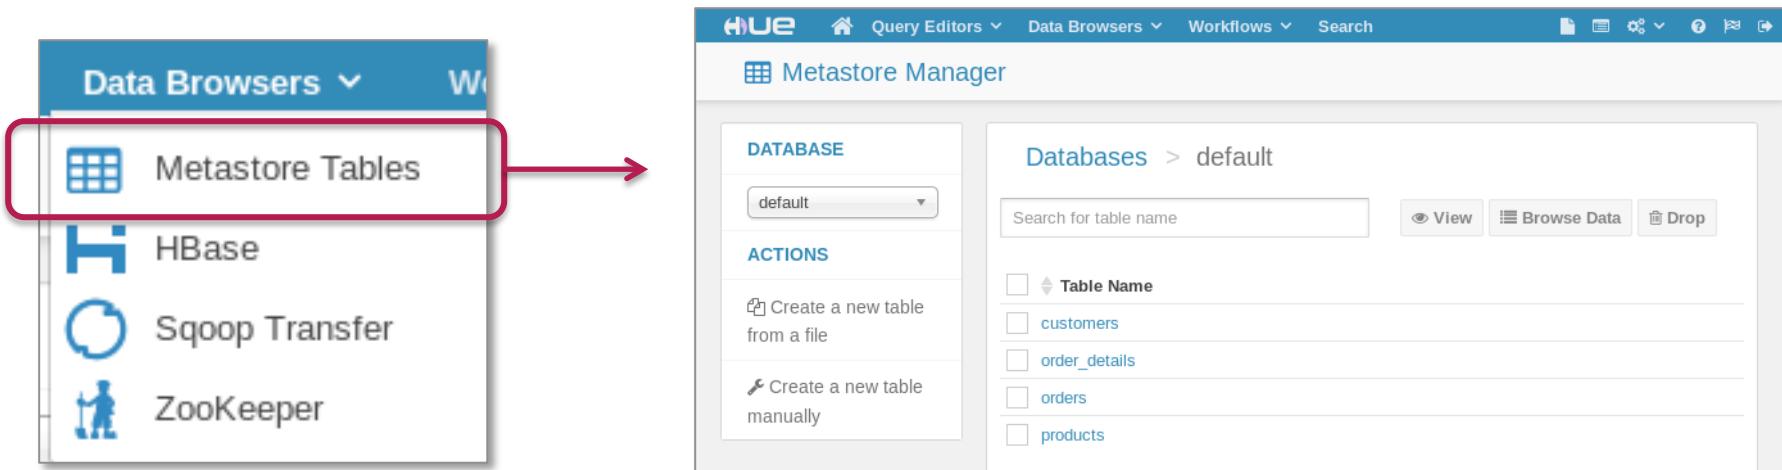
- **SHOW CREATE TABLE** displays the SQL command to create the table

```
SHOW CREATE TABLE jobs;
+-----+
| CREATE TABLE default.jobs
|   id INT,
|   title STRING,
|   salary INT,
|   posted TIMESTAMP
| )
| ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' '
...
...
```

# Using the Hue Metastore Manager

## ■ The Hue Metastore Manager

- An alternative to using SQL commands to manage metadata
- Allows you to create, load, preview, and delete databases and tables
  - Not all features are supported yet



# Chapter Topics

## Modeling and Managing Data With Impala and Hive

## Importing and Modeling Structured Data

- Data Storage Overview
- Creating Databases and Tables
- **Loading Data into Tables**
- HCatalog
- Impala Metadata Caching
- Conclusion
- Hands-On Exercise: Create and Populate Tables in Impala or Hive

## Data Validation

---

- **Impala and Hive are ‘schema on read’**
  - Unlike an RDBMS, they do not validate data on insert
    - Files are simply moved into place
    - Loading data into tables is therefore very fast
    - Errors in file format will be discovered when queries are performed
- **Missing or invalid data will be represented as NULL**

## Loading Data From HDFS Files

- To load data, simply add files to the table's directory in HDFS
  - Can be done directly using the `hdfs dfs` commands
  - This example loads data from HDFS into the `sales` table

```
$ hdfs dfs -mv \
 /tmp/sales.txt /user/hive/warehouse/sales/
```

- Alternatively, use the `LOAD DATA INPATH` command
  - Done from within Hive or Impala
  - This *moves* data within HDFS, just like the command above
  - Source can be either a file or directory

```
LOAD DATA INPATH '/tmp/sales.txt'
INTO TABLE sales;
```

## Overwriting Data From Files

---

- Add the **OVERWRITE** keyword to delete all records before import
  - Removes all files within the table's directory
  - Then moves the new files into that directory

```
LOAD DATA INPATH '/tmp/sales.txt'  
OVERWRITE INTO TABLE sales;
```

## Appending Selected Records to a Table

- Another way to populate a table is through a query
  - Use **INSERT INTO** to add results to an *existing* Hive table

```
INSERT INTO TABLE accounts_copy
SELECT * FROM accounts;
```

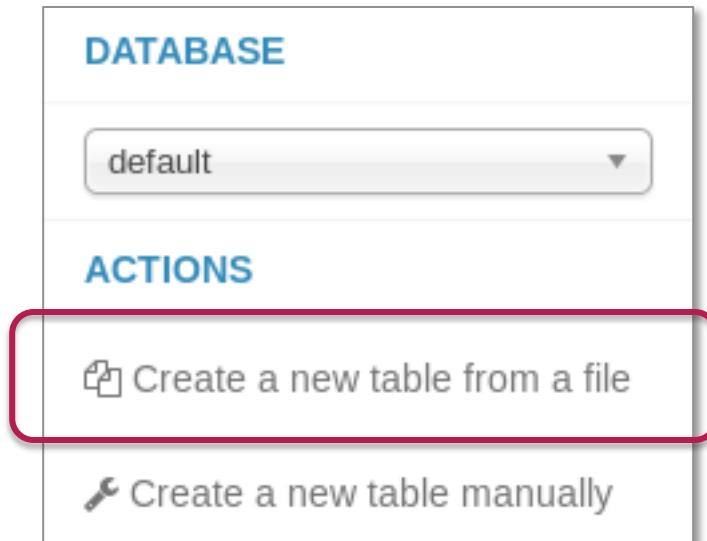
- Specify a **WHERE** clause to control which records are appended

```
INSERT INTO TABLE loyal_customers
SELECT * FROM accounts
WHERE YEAR(acct_create_dt) = 2008
AND acct_close_dt IS NULL;
```

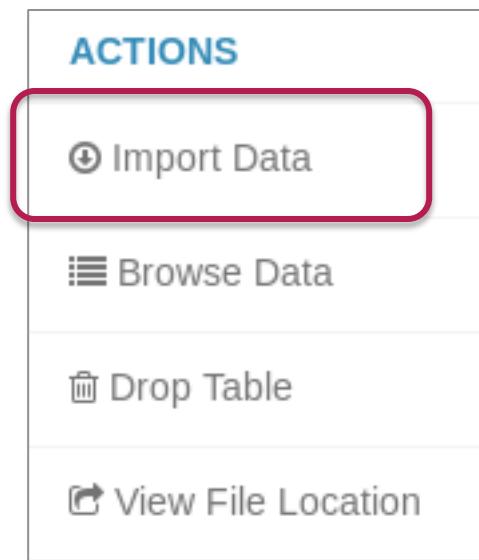
# Loading Data Using the Metastore Manager

- The Metastore Manager provides two ways to load data into a table

Table creation wizard



Import data wizard



## Loading Data From a Relational Database

- Sqoop has built-in support for importing data into Hive and Impala
- Add the **--hive-import** option to your Sqoop command
  - Creates the table in the Hive metastore
  - Imports data from the RDBMS to the table's directory in HDFS

```
$ sqoop import \
  --connect jdbc:mysql://localhost/loudacre \
  --username training \
  --password training \
  --fields-terminated-by '\t' \
  --table employees \
  --hive-import
```

- Note that **--hive-import** creates a table accessible in both Hive and Impala

# Chapter Topics

## Modeling and Managing Data With Impala and Hive

## Importing and Modeling Structured Data

- Data Storage Overview
- Creating Databases and Tables
- Loading Data into Tables
- **HCatalog**
- Impala Metadata Caching
- Conclusion
- Hands-On Exercise: Create and Populate Tables in Impala or Hive

## The Metastore and HCatalog

---

- **HCatalog is a Hive sub-project that provides access to the Metastore**
  - Accessible via command line and REST API
  - Allows you to define tables using HiveQL DDL syntax
  - Access those tables from Hive, Impala, MapReduce, Pig, and other tools
  - Included with CDH 4.2 and later

## Creating Tables in HCatalog

- HCatalog uses Hive's DDL (data definition language) syntax
  - You can specify a single command using the `-e` option

```
$ hcat -e "CREATE TABLE vendors \
(id INT, company STRING, email STRING) \
ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' \
LOCATION '/dualcore/vendors'"
```

- Tip: save longer commands to a text file and use the `-f` option
  - If the file has more than one command, separate each with a semicolon

```
$ hcat -f createtable.txt
```

## Displaying Metadata in HCatalog

- The SHOW TABLES command also shows tables created directly in Hive

```
$ hcat -e 'SHOW TABLES'  
employees  
vendors
```

- The DESCRIBE command lists the fields in a specified table
  - Use DESCRIBE FORMATTED instead to see detailed information

```
$ hcat -e 'DESCRIBE vendors'  
id      int  
company string  
email   string
```

## Removing a Table in HCatalog

---

- The **DROP TABLE** command has the same behavior as it does in Hive and Impala
  - Caution: this will remove the data as well as the metadata (unless table is **EXTERNAL**)!

```
$ hcat -e 'DROP TABLE vendors'
```

# Chapter Topics

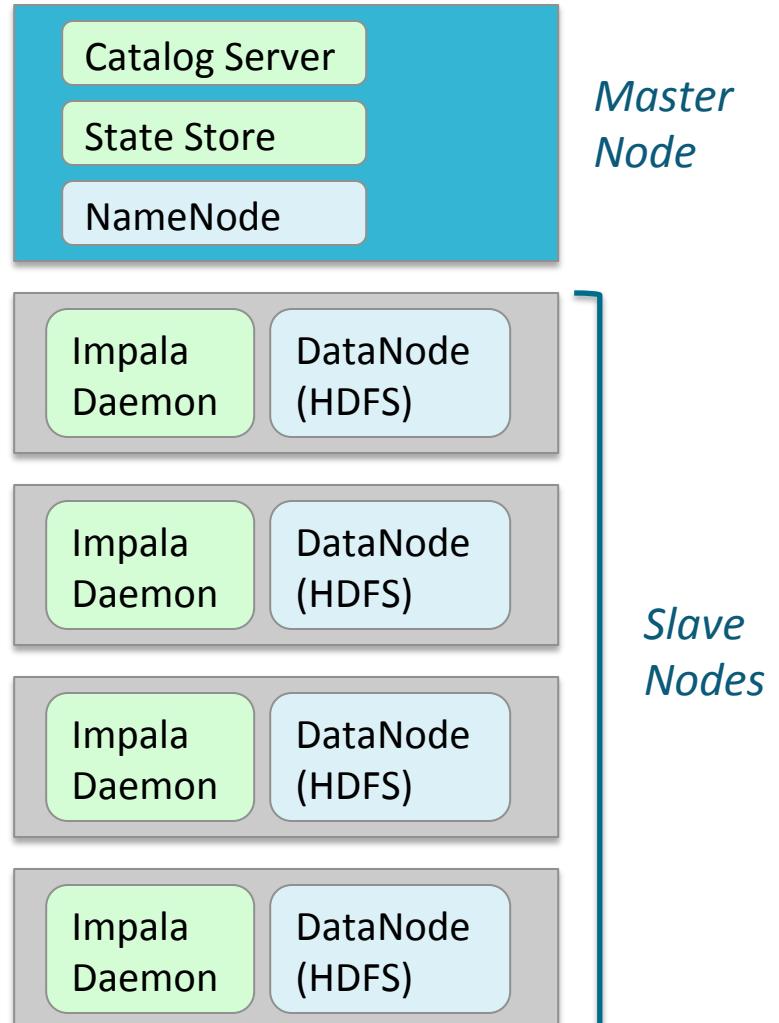
## Modeling and Managing Data With Impala and Hive

## Importing and Modeling Structured Data

- Data Storage Overview
- Creating Databases and Tables
- Loading Data into Tables
- HCatalog
- **Impala Metadata Caching**
- Conclusion
- Hands-On Exercise: Create and Populate Tables in Impala or Hive

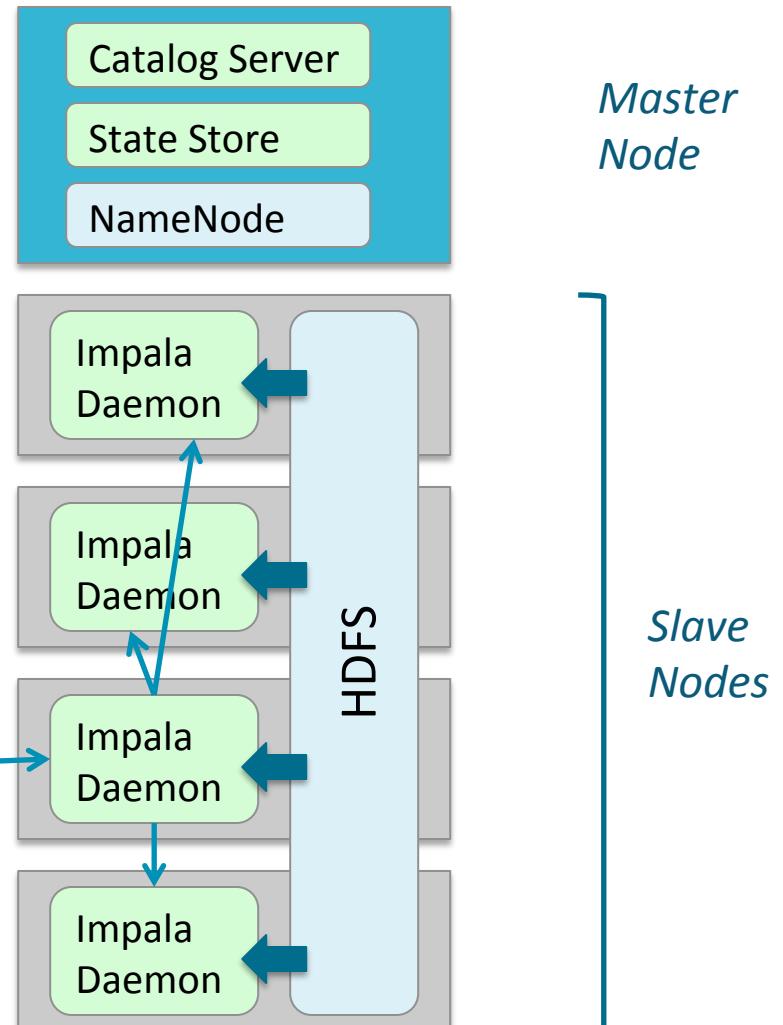
## Impala in the Cluster

- **Each slave node in the cluster runs an Impala daemon**
  - Co-located with the HDFS slave daemon (DataNode)
- **Two other daemons running on master nodes support query execution**
  - The **State Store** daemon
    - Provides lookup service for Impala daemons
    - Periodically checks status of Impala daemons
  - The **Catalog** daemon
    - Relays metadata changes to all the Impala daemons in a cluster



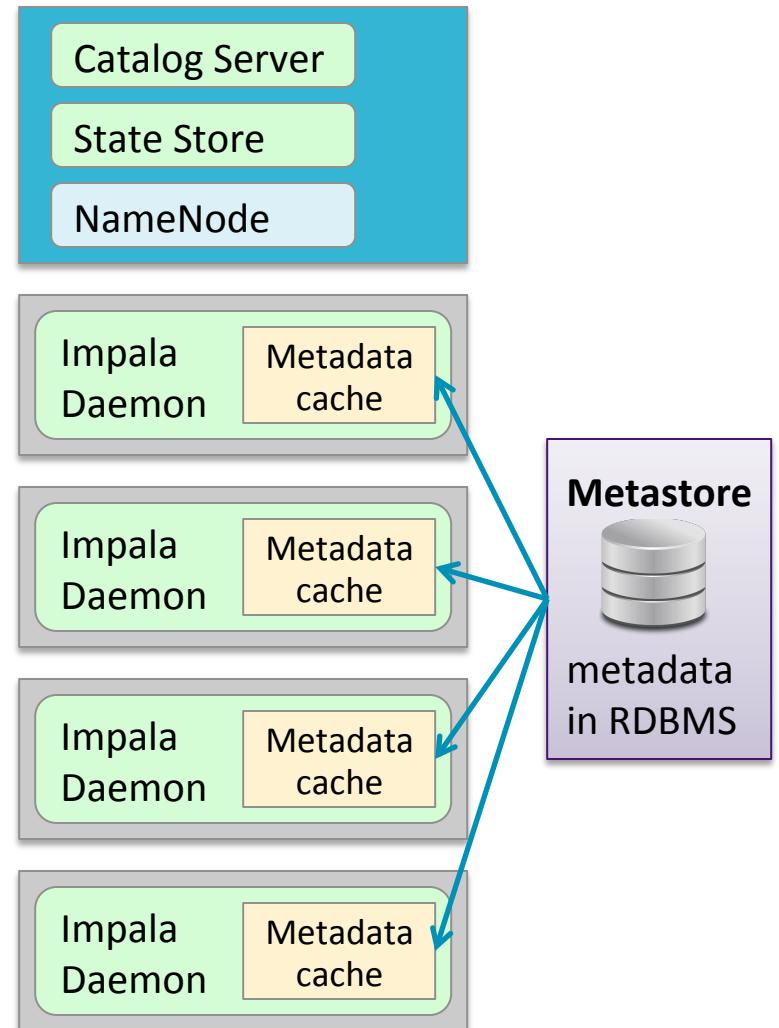
## How Impala Executes a Query

- **Impala daemon plans the query**
  - Client (impala-shell or Hue) connects to a local impala daemon
    - This is the *coordinator*
  - Coordinator requests a list of other Impala daemons in the cluster from the State Store
  - Coordinator distributes the query across other Impala daemons
  - Streams results to client



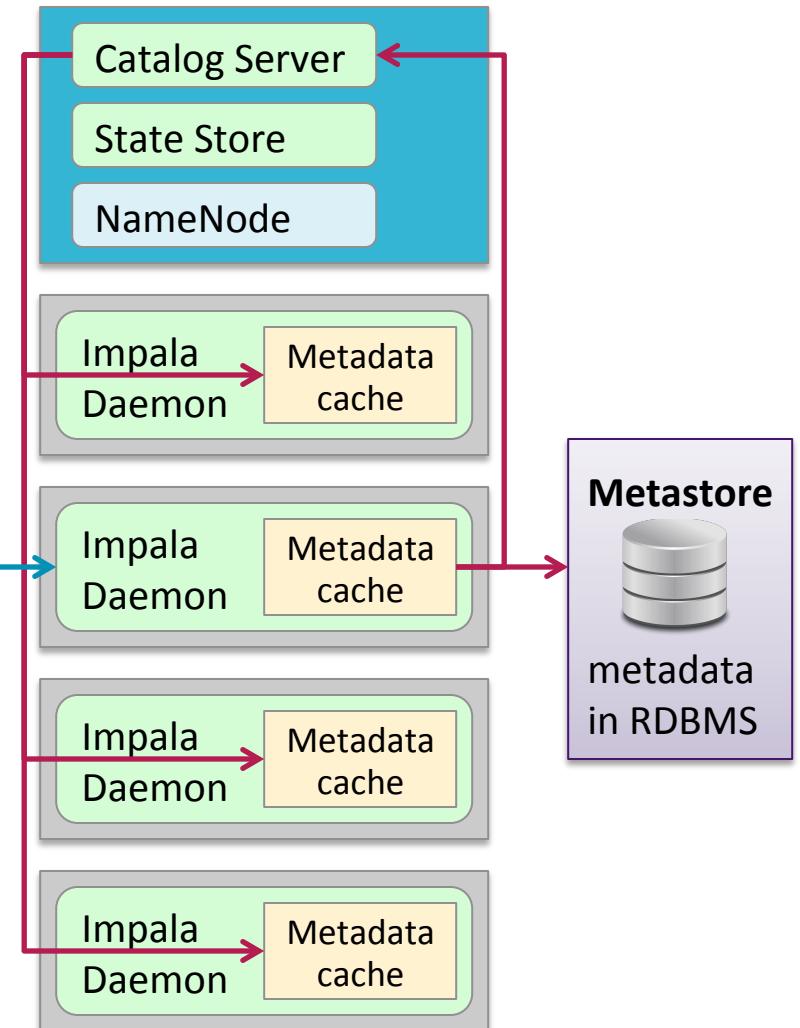
## Metadata Caching (1)

- **Impala daemons cache metadata**
  - The tables' schema definitions
  - The locations of tables' HDFS blocks
- **Metadata is cached from the Metastore at startup**



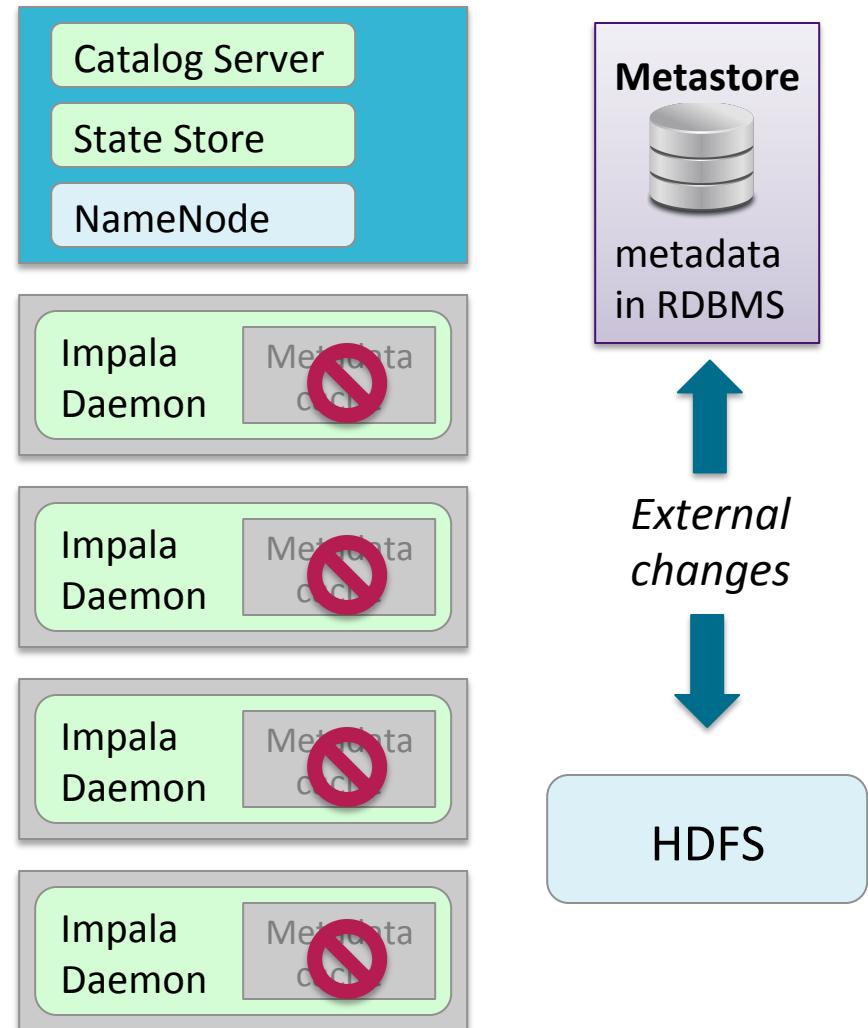
## Metadata Caching (2)

- When one Impala daemon changes the metastore, it notifies the catalog service
- The catalog service notifies all Impala daemons to update their cache



## External Changes and Metadata Caching

- **Metadata updates made *from outside of Impala* are not known to Impala, e.g.**
  - Changes via Hive, HCatalog or Hue Metadata Manager
  - Data added directly to directory in HDFS
- **Therefore the Impala metadata caches will be invalid**
- **You must manually refresh or invalidate Impala's metadata cache**



## Updating the Impala Metadata Cache

External Metadata Change	Required Action	Effect on Local Caches
New table added	<b>INVALIDATE METADATA</b> (with no table name)	Marks the entire cache as stale; metadata cache is reloaded as needed.
Table schema modified <i>or</i> New data added to a table	<b>REFRESH &lt;table&gt;</b>	Reloads the metadata for one table <i>immediately</i> . Reloads HDFS block locations for new data files only.
Data in a table extensively altered, such as by HDFS balancing	<b>INVALIDATE METADATA &lt;table&gt;</b>	Marks the metadata for a single table as stale. When the metadata is needed, all HDFS block locations are retrieved.

# Chapter Topics

## Modeling and Managing Data With Impala and Hive

## Importing and Modeling Structured Data

- Data Storage Overview
- Creating Databases and Tables
- Loading Data into Tables
- HCatalog
- Impala Metadata Caching
- **Conclusion**
- Hands-On Exercise: Create and Populate Tables in Impala or Hive

## Essential Points

---

- **Each table maps to a directory in HDFS**
  - Table data is stored as one or more files
  - Default format: plain text with delimited fields
- **The Metastore stores data *about* the data in an RDBMS**
  - E.g. Location, column names and types
- **Tables are created and managed using the Impala SQL or HiveQL Data Definition Language**
- **Impala caches metadata from the Metastore**
  - Invalidate or refresh the cache if tables are modified outside Impala
- **HCatalog provides access to the Metastore from tools outside Hive or Impala (e.g. Pig, MapReduce)**

## Bibliography

---

**The following offer more information on topics discussed in this chapter**

- **Impala Concepts and Architecture**
  - <http://tiny.cloudera.com/adcc12a>
- **Impala SQL Language Reference**
  - <http://tiny.cloudera.com/impalasql>
- **Impala-related Articles on Cloudera's Blog**
  - <http://tiny.cloudera.com/adcc12e>
- **Apache Hive Web Site**
  - <http://hive.apache.org/>
- **HiveQL Language Manual**
  - <http://tiny.cloudera.com/adcc10b>

# Chapter Topics

## Modeling and Managing Data With Impala and Hive

## Importing and Modeling Structured Data

- Data Storage Overview
- Creating Databases and Tables
- Loading Data into Tables
- HCatalog
- Impala Metadata Caching
- Conclusion
- **Hands-On Exercise: Create and Populate Tables in Impala or Hive**

## Hands-On Exercise: Create and Populate Tables in Impala

---

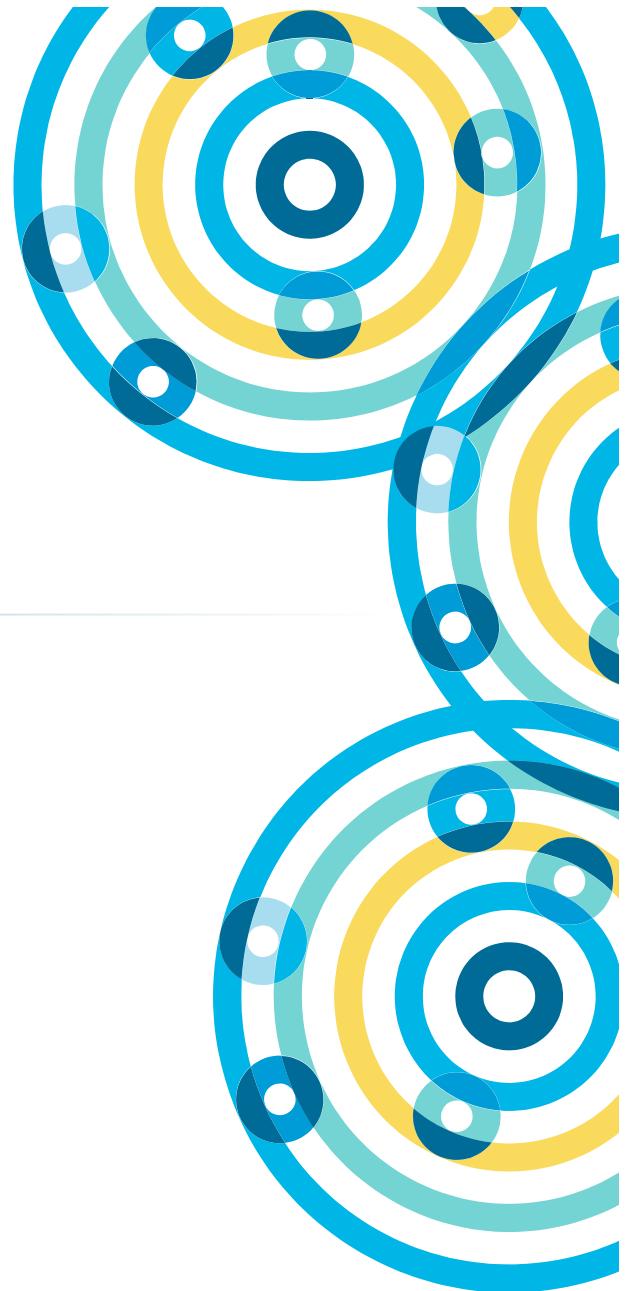
- **In this exercise you will**
  - Create a table in Impala to model and view existing data
  - Use Sqoop to create a new table automatically from data imported from MySQL
- **Please refer to the Hands-On Exercise Manual for instructions**



# Data Formats

---

## Chapter 7



# Course Chapters

- Introduction
- Introduction to Hadoop and the Hadoop Ecosystem
- Hadoop Architecture and HDFS
- Importing Relational Data with Apache Sqoop
- Introduction to Impala and Hive
- Modeling and Managing Data with Impala and Hive
- **Data Formats**
- Data File Partitioning
- Capturing Data with Apache Flume
- Spark Basics
- Working with RDDs in Spark
- Aggregating Data with Pair RDDs
- Writing and Deploying Spark Applications
- Parallel Processing in Spark
- Spark RDD Persistence
- Common Patterns in Spark Data Processing
- Spark SQL and DataFrames
- Conclusion

Course Introduction

Introduction to Hadoop

## Importing and Modeling Structured Data

Ingesting Streaming Data

Distributed Data Processing with Spark

Course Conclusion

## Data Formats

---

**In this chapter you will learn**

- **How to select the best data format for your needs**
- **How various Hadoop tools support different data formats**
- **How to define Avro schemas**
- **How Avro schemas can evolve to accommodate changing requirements**
- **How to extract data and metadata from an Avro data file**

# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- **File Formats**
  - Avro Schemas
  - Avro Schema Evolution
  - Using Avro with Impala, Hive and Sqoop
  - Using Parquet with Impala, Hive and Sqoop
  - Compression
  - Conclusion
  - Hands-On Exercise: Select a Format for a Data File

## File Storage Formats

---

- In previous chapters you saw that alternate file formats were available
  - E.g., in Hive and Impala

```
CREATE TABLE tablename (colname DATATYPE, . . .)
  ROW FORMAT DELIMITED
    FIELDS TERMINATED BY char
  STORED AS format
```

- E.g., in Sqoop
  - *--as-format*
- What formats are available?
- Which should you choose and why?

## Hadoop File Formats: Text Files

---

- **Text files are the most basic file type in Hadoop**
  - Can be read or written from virtually any programming language
  - Comma- and tab-delimited files are compatible with many applications
- **Text files are human readable, since everything is a string**
  - Useful when debugging
- **At scale, this format is inefficient**
  - Representing numeric values as strings wastes storage space
  - Difficult to represent binary data such as images
    - Often resort to techniques such as Base64 encoding
  - Conversion to/from native types adds performance penalty
- **Verdict: Good interoperability, but poor performance**

## Hadoop File Formats: Sequence Files

---

- **SequenceFiles store key-value pairs in a binary container format**
  - Less verbose and more efficient than text files
  - Capable of storing binary data such as images
  - Format is Java-specific and tightly coupled to Hadoop
- **Verdict: Good performance, but poor interoperability**

## Hadoop File Formats: Avro Data Files

---

- Efficient storage due to optimized binary encoding
- Widely supported throughout the Hadoop ecosystem
  - Can also be used outside of Hadoop
- Ideal for long-term storage of important data
  - Can read and write from many languages
  - Embeds schema in the file, so will always be readable
  - Schema evolution can accommodate changes
- Verdict: Excellent interoperability and performance
  - Best choice for general-purpose storage in Hadoop
- *More detail in coming slides*



## Columnar Formats

- Hadoop also supports **columnar format**
  - These organize data storage by column, rather than by row
  - Very efficient when selecting only a small subset of a table's columns

id	name	city	occupation	income	phone
1	Alice	Palo Alto	Accountant	85000	650-555-9748
2	Bob	Sunnyvale	Accountant	81500	650-555-8865
3	Bob	Palo Alto	Dentist	196000	650-555-7185
4	Bob	Palo Alto	Manager	87000	650-555-2518
5	Carol	Palo Alto	Manager	79000	650-555-3951
6	David	Sunnyvale	Mechanic	62000	650-555-4754

*Organization of data in traditional row-based formats*

id	name	city	occupation	income	phone
1	Alice	Palo Alto	Accountant	85000	650-555-9748
2	Bob	Sunnyvale	Accountant	81500	650-555-8865
3	Bob	Palo Alto	Dentist	196000	650-555-7185
4	Bob	Palo Alto	Manager	87000	650-555-2518
5	Carol	Palo Alto	Manager	79000	650-555-3951
6	David	Sunnyvale	Mechanic	62000	650-555-4754

*Organization of data in columnar formats*

## Hadoop File Formats: Parquet Files

---

- **Parquet is a columnar format developed by Cloudera and Twitter**
  - Supported in Spark, MapReduce, Hive, Pig, Impala, Crunch, and others
  - Schema metadata is embedded in the file (like Avro)
- **Uses advanced optimizations described in Google's Dremel paper**
  - Reduces storage space
  - Increases performance
- **Most efficient when adding many records at once**
  - Some optimizations rely on identifying repeated patterns
- **Verdict: Excellent interoperability and performance**
  - Best choice for column-based access patterns



# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- Selecting a File Format
- Hadoop Tool Support for File Formats
- **Avro Schemas**
- Using Avro with Hive and Sqoop
- Avro Schema Evolution
- Compression
- Conclusion
- Hands-On Exercise: Working with Avro Schemas

## Data Serialization

---

- **To understand Avro, you must first understand *serialization***
  - A way of representing data in memory as a series of bytes
  - Allows you to save data to disk or send it across the network
  - *Deserialization* allows you to read that data back into memory
- **For example, how do you serialize the number 108125150?**
  - 4 bytes when stored as a Java `int`
  - 9 bytes when stored as a Java `String`
- **Many programming languages and libraries support serialization**
  - Such as `Serializable` in Java or `pickle` in Python
- **Backwards compatibility and cross-language support can be challenging**
  - Avro was developed to address these challenges

## What is Apache Avro?

---

- **Avro Data File Format is just one part of the Avro project**
  - But it is the part this course focuses on
- **Avro is an efficient data serialization framework**
  - Top-level Apache project created by Doug Cutting (creator of Hadoop)
  - Widely supported throughout Hadoop and its ecosystem
- **Offers compatibility without sacrificing performance**
  - Data is serialized according to a *schema* you define
  - Read/write data in Java, C, C++, C#, Python, PHP, and other languages
  - Serializes data using a highly-optimized binary encoding
  - Specifies rules for *evolving* your schema over time
- **Avro also supports Remote Procedure Calls (RPC)**
  - Can be used for building custom network protocols
  - Flume uses this for internal communication

## Supported Types in Avro Schemas (Simple)

---

- A simple type holds exactly one value

Name	Description	Java Equivalent
<b>null</b>	An absence of a value	<b>null</b>
<b>boolean</b>	A binary value	<b>boolean</b>
<b>int</b>	32-bit signed integer	<b>int</b>
<b>long</b>	64-bit signed integer	<b>long</b>
<b>float</b>	Single-precision floating point value	<b>float</b>
<b>double</b>	Double-precision floating point value	<b>double</b>
<b>bytes</b>	Sequence of 8-bit unsigned bytes	<b>java.nio.ByteBuffer</b>
<b>string</b>	Sequence of Unicode characters	<b>java.lang.CharSequence</b>

## Supported Types in Avro Schemas (Complex)

---

- Avro also supports complex types

Name	Description
<b>record</b>	A user-defined type composed of one or more named fields
<b>enum</b>	A specified set of values
<b>array</b>	Zero or more values of the same type
<b>map</b>	Set of key-value pairs; key is string while value is of specified type
<b>union</b>	Exactly one value matching a specified set of types
<b>fixed</b>	A fixed number of 8-bit unsigned bytes

- The **record** type is the most important
  - Main use of other types is to define a record's fields

## Basic Schema Example

- Excerpt from a SQL CREATE TABLE statement

```
CREATE TABLE employees
  (id INT, name STRING, title STRING, bonus INT)
```

- Equivalent Avro schema

```
{"namespace": "com.loudacre.data",
"type": "record",
"name": "Employee",
"fields": [
    {"name": "id", "type": "int"},
    {"name": "name", "type": "string"},
    {"name": "title", "type": "string"},
    {"name": "bonus", "type": "int"}]
```

## Specifying Default Values in the Schema

- Avro also supports setting a default value in the schema
  - Used when no value was explicitly set for a field
  - Similar to SQL

```
{"namespace": "com.loudacre.data",
"type": "record",
"name": "Invoice",
"fields": [
    {"name": "id", "type": "int"},
    {"name": "taxcode", "type": "int", "default": "39"},
    {"name": "lang", "type": "string", "default": "EN_US"}]
```

The `taxcode` and `lang` fields have default values

## Avro Schemas and Null Values

- Avro checks for null values when serializing the data
- Null values are only allowed *when explicitly specified* in the schema

```
{"namespace": "com.loudacre.data",
 "type": "record",
 "name": "Employee",
 "fields": [
     {"name": "id", "type": "int"},
     {"name": "name", "type": "string"},
     {"name": "title", "type": ["null", "string"]},
     {"name": "bonus", "type": ["null", "int"]}]
}
```

The **title** and **bonus** fields allow null values

## Schema Example with Complex Types

- The following example shows a record with an enum and a string array

```
{ "namespace" : "com.loudacre.data",
  "type" : "record",
  "name" : "CustomerServiceTicket",
  "fields" : [
    { "name" : "id", "type" : "int" },
    { "name" : "agent", "type" : "string" },
    { "name" : "category", "type" : {
        "name" : "CSCategory", "type" : "enum",
        "symbols" : [ "Order", "Shipping", "Device" ] }
    },
    { "name" : "tags", "type" : {
        "type" : "array", "items" : "string" }
    }
}
```

The **category** field has three enumerated possible values

**tags** is an array of strings

## Documenting Your Schema

- It's a good practice to document any ambiguities in a schema
  - All types (including record) support an optional doc attribute

```
{ "namespace": "com.loudacre.data",
  "type": "record",
  "name": "WebProduct",
  "doc": "Item currently sold in Loudacre's online store",
  "fields": [
    { "name": "id", "type": "int", "doc": "Product SKU" },
    { "name": "shipwt", "type": "int",
      "doc": "Shipping weight, in pounds" },
    { "name": "price", "type": "int",
      "doc": "Retail price, in cents (US) " }
  ]
}
```

## Avro Container Format

---

- **Avro also defines a container file format for storing Avro records**
  - Also known as “Avro data file format”
  - Similar to Hadoop SequenceFile format
  - Cross-language support for reading and writing data
- **Supports compressing *blocks* (groups) of records**
  - It is “splittable” for efficient processing in Hadoop
- **This format is self-describing**
  - Each file contains a copy of the schema used to write its data
  - All records in a file must use the same schema

# Inspecting Avro Data Files with Avro Tools

- **Avro data files are an efficient way to store data**
  - However, the binary format makes debugging difficult
- **Each Avro release contains an Avro Tools JAR file**
  - Allows you to read the schema or data for an Avro file
  - Included with CDH 5 and later
  - Available for download from the Avro Web site or Maven repository

```
$ java -jar avro-tools*.jar tojson mydatafile.avro
{"name": "Alice", "salary": 56500, "city": "Anaheim"}
 {"name": "Bob", "salary": 51400, "city": "Bellevue"}

$ java -jar avro-tools*.jar getschema mydatafile.avro
{
  "type" : "record",
  "name" : "DeviceData",
  "namespace" : "com.loudacre.data", ...rest of schema follows
```

# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- Selecting a File Format
- Hadoop Tool Support for File Formats
- Avro Schemas
- Using Avro with Impala, Hive and Sqoop
- **Avro Schema Evolution**
- Compression
- Conclusion
- Hands-On Exercise: Select a Format for a Data File

## Schema Evolution

---

- **The structure of your data will change over time**
  - Fields may be added, removed, changed, or renamed
  - In SQL, these are handled with **ALTER TABLE** statements
- **These changes can break compatibility with many formats**
  - Objects serialized in **SequenceFiles** become unreadable
- **Data written to Avro data files is always readable**
  - The schema used to write the data is embedded in the file itself
  - However, an application reading data might expect the *new* structure
- **Avro has a unique approach to maintaining forward compatibility**
  - A reader can use a different schema than the writer

## Schema Evolution: A Practical Example (1)

- Imagine that we have written millions of records with this schema

```
{ "namespace": "com.loudacre.data",
  "type": "record",
  "name": "CustomerContact",
  "fields": [
    { "name": "id", "type": "int" },
    { "name": "name", "type": "string" },
    { "name": "faxNumber", "type": "string" }
  ]
}
```

## Schema Evolution: A Practical Example (2)

- We would like to modernize this based on the schema below
  - Rename id field to customerId and change type from int to long
  - Remove faxNumber field
  - Add prefLang field
  - Add email field

```
{ "namespace": "com.loudacre.data",
  "type": "record",
  "name": "CustomerContact",
  "fields": [
    {"name": "customerId", "type": "long"},
    {"name": "name", "type": "string"},
    {"name": "prefLang", "type": "string"},
    {"name": "email", "type": "string"}
  ]
}
```

## Schema Evolution: A Practical Example (3)

---

- **We could use the new schema to write new data**
  - Applications that use the new schema could read the new data
- **Unfortunately, new applications wouldn't be able to read the *old* data**
  - We must make a few schema changes to improve compatibility

## Schema Evolution: A Practical Example (4)

- If you rename a field, you must specify an alias for the old name(s)
  - Here, we map the old `id` field to the new `customerId` field

```
{ "namespace": "com.loudacre.data",
  "type": "record",
  "name": "CustomerContact",
  "fields": [
    { "name": "customerId", "type": "long",
      "aliases": ["id"] },
    { "name": "name", "type": "string" },
    { "name": "prefLang", "type": "string" },
    { "name": "email", "type": "string" }
  ]
}
```

## Schema Evolution: A Practical Example (5)

- Newly-added fields will lack values for records previously written
  - You must specify a default value

```
{ "namespace": "com.loudacre.data",
  "type": "record",
  "name": "CustomerContact",
  "fields": [
    { "name": "customerId", "type": "long",
      "aliases": ["id"] },
    { "name": "name", "type": "string" },
    { "name": "prefLang", "type": "string",
      "default": "en_US" },
    { "name": "email",
      "type": ["null", "string"], "default": null }
  ]
}
```

Default value for  
prefLang is  
`en_US`

`email` is nullable  
so `null` can be the  
default

## Schema Evolution: Compatible Changes

---

- **The following changes will not affect existing readers**
  - Adding, changing, or removing a `doc` attribute
  - Changing a field's default value
  - Adding a new field with a default value
  - Removing a field that specified a default value
  - Promoting a field to a wider type (e.g., `int` to `long`)
  - Adding aliases for a field

## Schema Evolution: Incompatible Changes

---

- **The following are some changes that might break compatibility**
  - Changing the record's name or namespace attributes
  - Adding a new field without a default value
  - Removing a symbol from an enum
  - Removing a type from a union
  - Modifying a field's type to one that could result in truncation
- **To handle these incompatibilities**
  1. Read your old data (using the original schema)
  2. Modify data as needed in your application
  3. Write the new data (using the new schema)
- **Existing readers/writers may need to be updated to use new schema**

# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- Selecting a File Format
- Hadoop Tool Support for File Formats
- Avro Schemas
- Avro Schema Evolution
- **Using Avro with Impala, Hive and Sqoop**
- Compression
- Conclusion
- Hands-On Exercise: Select a Format for a Data File

## Using Avro with Sqoop

---

- Sqoop supports importing data as Avro, or exporting data from existing Avro data files
- `sqoop import` saves the schema JSON file in local directory

```
$ sqoop import \
  --connect jdbc:mysql://localhost/loudacre \
  --username training --password training \
  --table accounts \
  --target-dir /loudacre/accounts_avro \
  --as-avrodatafile
```

## Using Avro with Impala and Hive (1)

---

- **Hive and Impala support Avro**
  - Hive supports all Avro types
  - Impala does not support complex types
    - **enum, array**, etc.

## Using Avro with Impala and Hive (2)

- Table creation can include schema inline or in a separate file

```
CREATE TABLE order_details_avro
  STORED AS AVRO
  TBLPROPERTIES ('avro.schema.url'=
    'hdfs://localhost/loudacre/accounts_schema.json');
```

```
CREATE TABLE order_details_avro
  STORED AS AVRO
  TBLPROPERTIES ('avro.schema.literal'=
    ' { "name": "order",
      "type": "record",
      "fields": [
        { "name": "order_id", "type": "int" },
        { "name": "cust_id", "type": "int" },
        { "name": "order_date", "type": "string" }
      ] }');
```

# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- Selecting a File Format
- Avro Schemas
- Avro Schema Evolution
- Using Avro with Impala, Hive and Sqoop
- **Using Parquet with Impala, Hive and Sqoop**
- Compression
- Conclusion
- Hands-On Exercise: Select a Format for a Data File

## Using Parquet with Swoop

- Swoop supports importing data as Parquet, or exporting data from existing Parquet data files
  - Use the **--as-parquetfile** option

```
$ swoop import \
--connect jdbc:mysql://localhost/loudacre \
--username training --password training \
--table accounts \
--target-dir /loudacre/accounts_parquet \
--as-parquetfile
```

## Using Parquet with Hive and Impala (1)

- Create a new table stored in Parquet format

```
CREATE TABLE order_details_parquet (
    order_id INT,
    prod_id INT)
STORED AS PARQUET;
```

\* **STORED AS PARQUET** supported in Impala, and in Hive 0.13 and later

## Using Parquet with Hive and Impala (2)

- In Impala, use **LIKE PARQUET** to use column metadata from an existing Parquet data file
- Example: Create a new table to access existing Parquet format data



```
CREATE EXTERNAL TABLE ad_data
  LIKE PARQUET '/loudacre/ad_data/datafile1.parquet'
  STORED AS PARQUET
  LOCATION '/loudacre/ad_data/' ;
```

# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- Selecting a File Format
- Avro Schemas
- Avro Schema Evolution
- Using Avro with Impala, Hive and Sqoop
- Using Parquet with Impala, Hive and Sqoop
- **Compression**
- Conclusion
- Hands-On Exercise: Select a Format for a Data File

## Performance Considerations

---

- You have just learned how different file formats affect performance
- Another factor can significantly affect performance: data compression

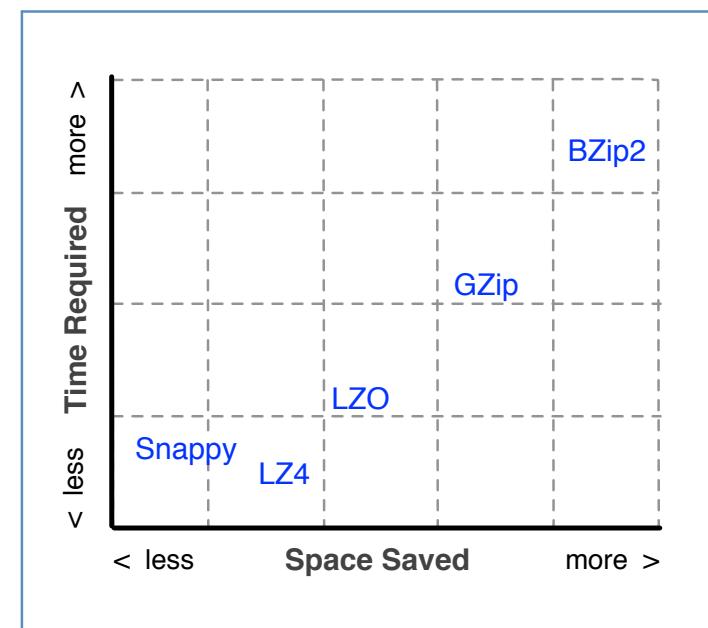
## Data Compression

- **Each file format may also support compression**
  - This reduces amount of disk space required to store data
- **Compression is a tradeoff between CPU time and bandwidth/storage space**
  - Aggressive algorithms take a long time, but save more space
  - Less aggressive algorithms save less space but are much faster
- **Can significantly improve performance**
  - Many Hadoop jobs are I/O-bound
  - Using compression allows you to handle more data per I/O operation
  - Compression can also improve the performance of network transfers



## Compression Codecs

- The implementation of a compression algorithm is known as a *codec*
  - Short for compressor/decompressor
- Many codecs are commonly used with Hadoop
  - Each has different performance characteristics
  - Not all Hadoop tools are compatible with all codecs
- Overall, BZip2 saves the most space
  - But LZ4 and Snappy are much faster
  - Impala supports Snappy but not LZ4
- For "hot" data, speed matters most
  - Better to compress by 40% in one second than by 80% in 10 seconds



## Using Compression With Sqoop

---

- Sqoop – use **--compression-codec** flag

- Example

```
--compression-codec  
org.apache.hadoop.io.compress.SnappyCodec
```

# Using Compression With Impala and Hive

- Not all file format/compression combinations are supported
  - Properties and syntax varies
  - See the documentation for a full list of supported formats and codecs for Impala and Hive
  - Caution: Impala queries data in memory – both compressed and uncompressed data are stored in memory
- Impala example

```
> CREATE TABLE mytable_parquet LIKE mytable_text  
      STORED AS PARQUET;  
> set PARQUET_COMPRESSION_CODEC=snappy;  
> INSERT INTO mytable_parquet  
      SELECT * FROM mytable_text;
```

# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- Selecting a File Format
- Hadoop Tool Support for File Formats
- Avro Schemas
- Using Avro with Hive and Sqoop
- Avro Schema Evolution
- Compression
- **Conclusion**
- Hands-On Exercise: Select a Format for a Data File

## Essential Points (1)

---

- **Hadoop and its ecosystem support many file formats**
  - May ingest data in one format, but convert to another as needed
- **Selecting the format for your data set involves several considerations**
  - Ingest pattern
  - Tool compatibility
  - Expected lifetime
  - Storage and performance requirements

## Essential Points (2)

---

- **Choose from the three main Hadoop file format options**
  - Text – Good for testing and interoperability
  - Avro – Best for general purpose performance and evolving schemas
  - Parquet – Best performance for column-oriented access patterns
- **Avro is a serialization framework that includes a data file format**
  - Compact binary encodings provide good performance
  - Supports schema evolution for long-term storage
- **Compression saves disk storage space and IO times at the cost of CPU time**
  - Hadoop tools support a number of different codecs

## Bibliography

---

The following offer more information on topics discussed in this chapter

- Avro Getting Started Guide (Java)
  - <http://tiny.cloudera.com/adcc03a>
- Avro Specification
  - <http://tiny.cloudera.com/adcc03b>
- Parquet
  - <https://parquet.apache.org>
- Announcing Parquet 1.0: Columnar Storage for Hadoop
  - <http://tiny.cloudera.com/adcc03c>

# Chapter Topics

## Data Formats

## Importing and Modeling Structured Data

- Selecting a File Format
- Avro Schemas
- Avro Schema Evolution
- Using Avro with Impala, Hive and Sqoop
- Using Parquet with Impala, Hive and Sqoop
- Compression
- Conclusion
- **Hands-On Exercise: Select a Format for a Data File**

## Hands-On Exercise: Select a Format for a Data File

---

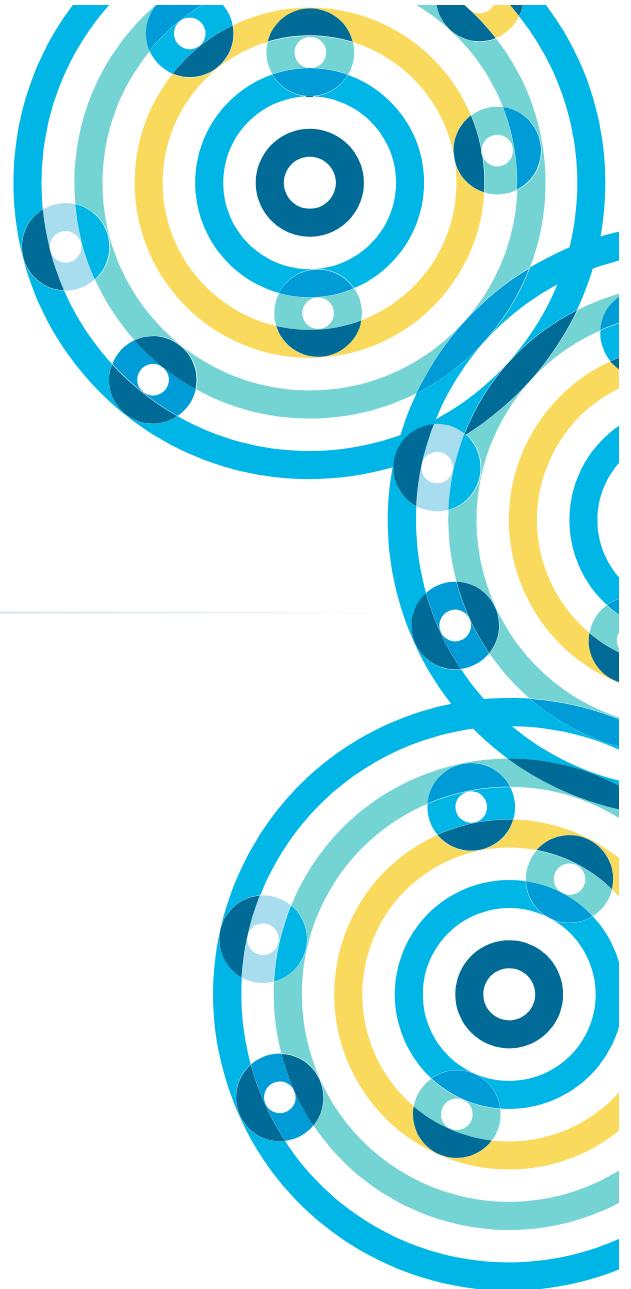
- **In this exercise you will**
  - Use Sqoop to import the accounts table in Avro format
  - Define an Impala table to access the Avro accounts data
  - Bonus: Save an existing plain text Impala table as Parquet
- **Please refer to the Hands-On Exercise Manual for instructions**



# Data File Partitioning

---

Chapter 8



# Course Chapters

- Introduction
- Introduction to Hadoop and the Hadoop Ecosystem
- Hadoop Architecture and HDFS
- Importing Relational Data with Apache Sqoop
- Introduction to Impala and Hive
- Modeling and Managing Data with Impala and Hive
- Data Formats
- **Data File Partitioning**
- Capturing Data with Apache Flume
- Spark Basics
- Working with RDDs in Spark
- Aggregating Data with Pair RDDs
- Writing and Deploying Spark Applications
- Parallel Processing in Spark
- Spark RDD Persistence
- Common Patterns in Spark Data Processing
- Spark SQL and DataFrames
- Conclusion

Course Introduction

Introduction to Hadoop

## Importing and Modeling Structured Data

Ingesting Streaming Data

Distributed Data Processing with Spark

Course Conclusion

# Data File Partitioning

---

**In this chapter you will learn**

- **How to improve query performance with data file partitioning**
- **How to create and populate partitioned tables in Impala and Hive**

# Chapter Topics

## Data File Partitioning

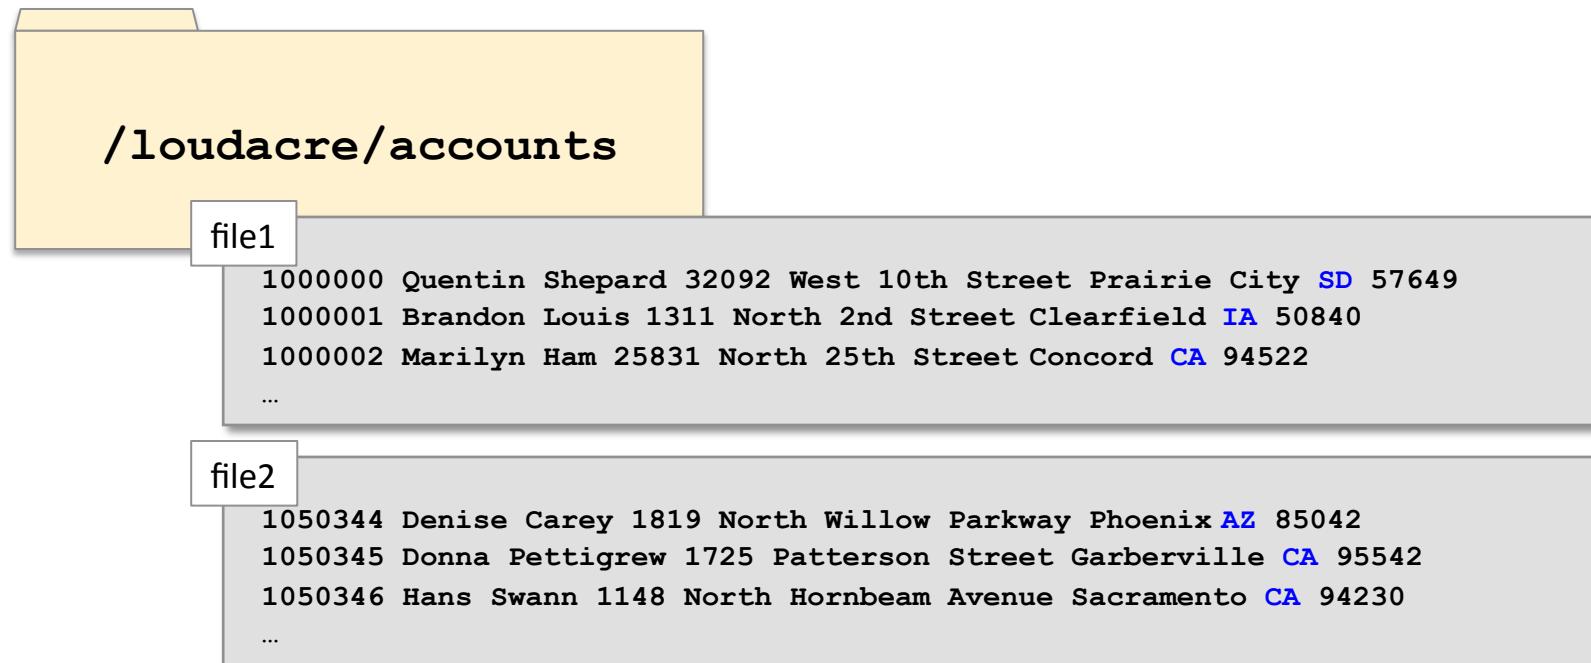
## Importing and Modeling Structured Data

### ■ Partitioning Overview

- Partitioning in Impala and Hive
- Conclusion
- Hands-On Exercise: Partition Data in Impala or Hive

## Data Storage Partitioning (1)

- By default, all files in a data set are stored in a single HDFS directory
  - All files in the directory are read during analysis or processing
  - “Full table scan”

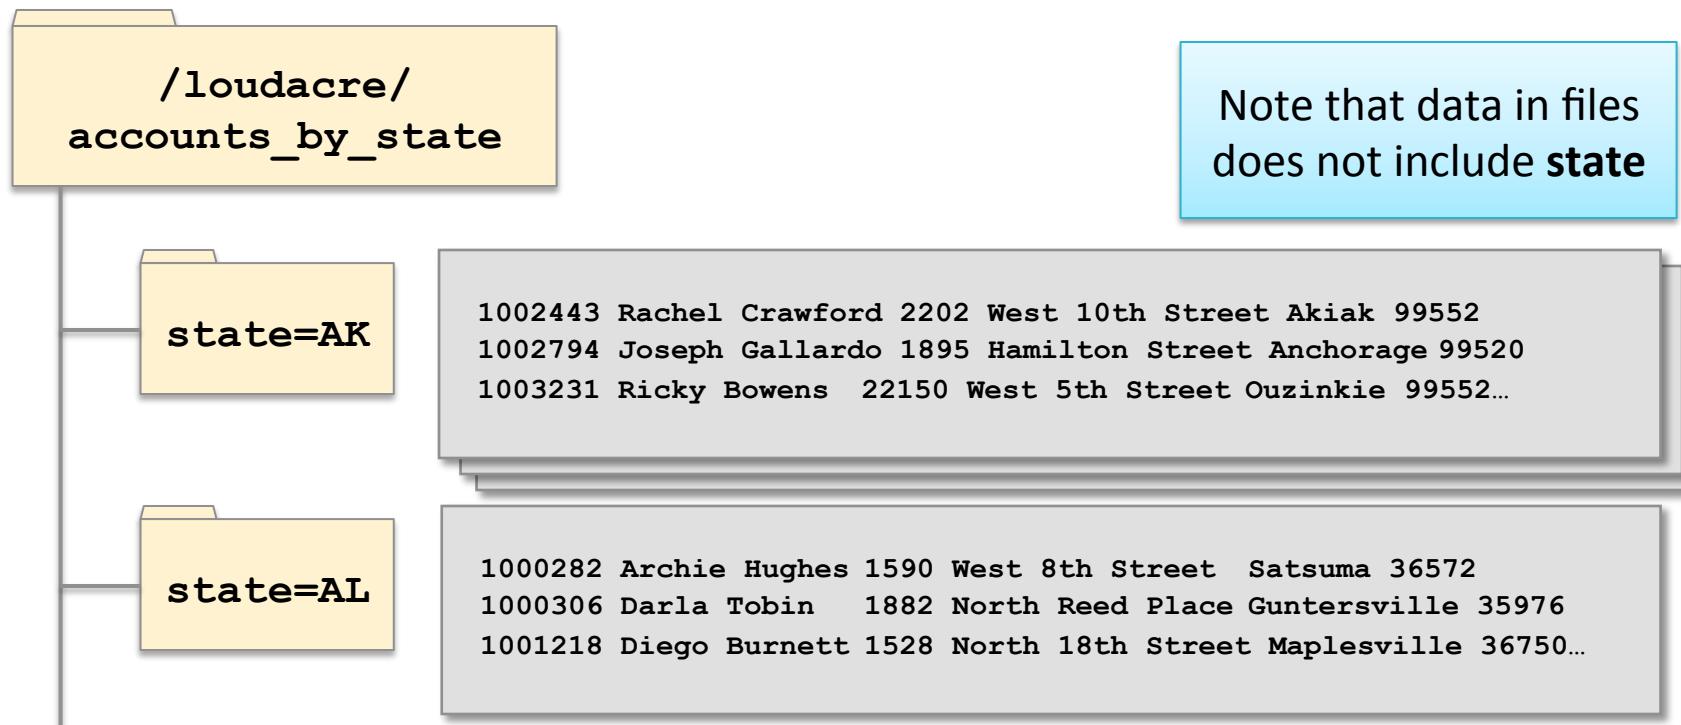


## Data Storage Partitioning (2)

- **Partitioning** subdivides the data

- Analysis can be done on only the relevant subset of data
  - Potentially much faster!

- Hadoop partitions using subdirectories



# Hadoop Partitioning

---

- **Partitioning is involved at two phases**
  - Storage – putting the data into correct partition (subdirectory)
  - Retrieval – getting the data out of the correct partition based on the query or analysis being done
- **Hadoop with built-in support for partitioning**
  - Hive and Impala (covered in next section)
  - Sqoop – When using the `--hive-import` option you can specify flags `--hive-partition-key` and `--hive-partition-value`
- **Other tools can be used to store partitioned data**
  - Spark and MapReduce
  - Flume (at ingestion)

# Chapter Topics

## Data File Partitioning

## Importing and Modeling Structured Data

- Partitioning Overview
- **Partitioning in Impala and Hive**
- Conclusion
- Hands-On Exercise: Partition Data in Impala or Hive

## Example: Impala/Hive Partitioning Accounts By State (1)

- Example: accounts is a non-partitioned table

```
CREATE EXTERNAL TABLE accounts(
    cust_id INT,
    fname STRING,
    lname STRING,
    address STRING,
    city STRING,
    state STRING,
    zipcode STRING)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/loudacre/accounts';
```

## Example: Impala/Hive Partitioning Accounts By State (2)

- What if most of Loudacre's analysis on the customer table was done by state? For example:

```
SELECT fname, lname  
      FROM accounts  
    WHERE state='NY';
```

- By default, all queries have to scan *all* files in the directory
- Use partitioning to store data in separate files by state
  - State-based queries scan only the relevant files

## Example: Impala/Hive Partitioning Accounts By State (3)

- Create a partitioned table using PARTITIONED BY

```
CREATE EXTERNAL TABLE accounts_by_state(
    cust_id INT,
    fname STRING,
    lname STRING,
    address STRING,
    city STRING,
    state STRING,
    zipcode STRING)
PARTITIONED BY (state STRING)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/loudacre/accounts_by_state';
```

## Partition Columns

- The partition column is displayed if you DESCRIBE the table

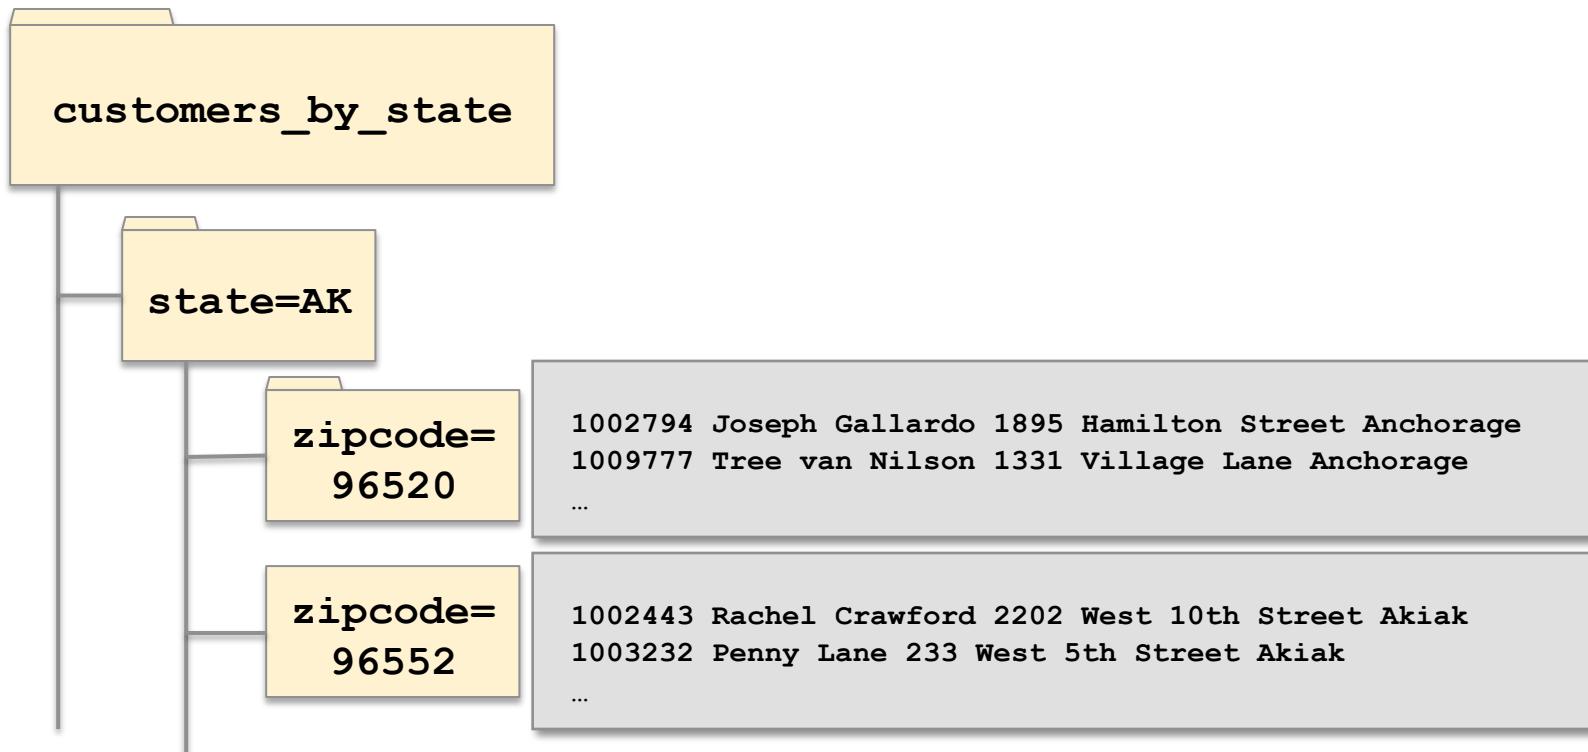
```
DESCRIBE accounts_by_state;
+-----+-----+-----+
| name      | type      | comment |
+-----+-----+-----+
| cust_id   | int       |          |
| fname     | string    |          |
| lname     | string    |          |
| address   | string    |          |
| city      | string    |          |
| zipcode   | string    |          |
| state     | string    |          |
+-----+-----+-----+
```

A partition column is a “virtual column”; data is not stored in the file

## Nested Partitions

- You can also create nested partitions

```
... PARTITIONED BY (state STRING, zipcode STRING)
```



## Loading Data Into a Partitioned Table

---

- **Dynamic partitioning**
  - Impala/Hive add new partitions automatically as needed at load time
  - Data is stored into the correct partition (subdirectory) based on column value
- **Static partitioning**
  - You define new partitions using ADD PARTITION
  - When loading data, you specify which partition to store data in

## Dynamic Partitioning

- We can create new partitions dynamically from existing data

```
INSERT OVERWRITE TABLE accounts_by_state
PARTITION(state)
SELECT cust_id, fname, lname, address,
city, zipcode, state FROM accounts;
```

- Partitions are automatically created based on the value of the *last* column
  - If the partition does not already exist, it will be created
  - If the partition does exist, it will be overwritten

## Static Partitioning Example: Partition Calls by Day (1)

---

- Loudacre's customer service phone system generates logs detailing calls received
  - Analysts use this data to summarize previous days' calls
  - For example:

```
SELECT event_type, COUNT(event_type)
  FROM call_log
 WHERE call_date = '2014-10-01'
 GROUP BY event_type;
```

## Static Partitioning Example: Partition Calls by Day (2)

- Logs are generated daily, e.g.

call-20141001.log

```
19:45:19,312-555-7834,CALL_RECEIVED  
19:45:23,312-555-7834,OPTION_SELECTED,Shipping  
19:46:23,312-555-7834,ON_HOLD  
19:47:51,312-555-7834,AGENT_ANSWER,Agent ID N7501  
19:48:37,312-555-7834,COMPLAINT,Item not received  
19:48:41,312-555-7834,CALL_END,Duration: 3:22
```

...

call-20141002.log

```
03:45:01,505-555-2345,CALL_RECEIVED  
03:45:09,505-555-2345,OPTION_SELECTED,Billing  
03:56:21,505-555-2345,AGENT_ANSWER,Agent ID A1503  
03:57:01,505-555-2345,QUESTION
```

...

## Static Partitioning Example: Partition Calls by Day (3)

---

- In the previous example, existing data was partitioned dynamically based on a column value
- This time we use static partitioning
  - Because the data files do not include the partitioning data

## Static Partitioning Example: Partition Calls by Day (4)

- The partitioned table is defined the same way

```
CREATE TABLE call_logs (
    call_time STRING,
    phone STRING,
    event_type STRING,
    details STRING)
PARTITIONED BY (call_date STRING)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';
```

## Loading Data Into Static Partitions (1)

- With static partitioning, you create new partitions as needed
- e.g. For each new day of call log data, add a partition:

```
ALTER TABLE call_logs
ADD PARTITION (call_date='2014-10-02');
```

- This command
  1. Adds the partition to the table's metadata
  2. Creates subdirectory  
/user/hive/warehouse/call\_logs/  
call\_date=2014-10-02

## Loading Data Into Static Partitions (2)

- Then load the day's data into the correct partition

```
LOAD DATA INPATH '/mystaging/call-20141002.log'  
    INTO TABLE call_logs  
    PARTITION(call_date='2014-10-02');
```

- This command moves the HDFS file `call-20141002.log` to the partition subdirectory
- To overwrite all data in a partition

```
LOAD DATA INPATH '/mystaging/call-20141002.log'  
    INTO TABLE call_logs OVERWRITE  
    PARTITION(call_date='2014-10-02');
```

## Hive Only: Shortcut for Loading Data Into Partitions

- Hive will create a new partition if the one specified doesn't exist



```
LOAD DATA INPATH '/mystaging/call-20141002.log'  
    INTO TABLE call_logs  
    PARTITION(call_date='2014-10-02');
```

- This command

1. Adds the partition to the table's metadata if it doesn't exist
2. Creates subdirectory  
`/user/hive/warehouse/call_logs/call_date=2014-10-02` if it doesn't exist
3. Moves the HDFS file `call-20141002.log` to the partition subdirectory

## Viewing, Adding, and Removing Partitions

- To view the current partitions in a table

```
SHOW PARTITIONS call_logs;
```

- Use ALTER TABLE to add or drop partitions

```
ALTER TABLE call_logs
  ADD PARTITION (call_date='2013-06-05')
  LOCATION '/loudacre/call_logs/call_date=2013-06-05' ;
```

```
ALTER TABLE call_logs
  DROP PARTITION (call_date='2013-06-06') ;
```

## Creating Partitions from Existing Partition Directories in HDFS

- Partition directories in HDFS can be created and populated outside Hive or Impala
  - For example, by a Spark or MapReduce application
- In Hive, use the `MSCK REPAIR TABLE` command to create (or recreate) partitions for an existing table



```
MSCK REPAIR TABLE call_logs;
```

## When To Use Partitioning

---

- **Use partitioning for tables when**

- Reading the entire data set takes too long
  - Queries almost always filter on the partition columns
  - There are a reasonable number of different values for partition columns
  - The data generation or ETL process segments data by file or directory names
    - Partition column values are not in the data itself

## When Not To Use Partitioning

---

- **Avoid partitioning data into numerous small data files**
  - Don't partition on columns with too many unique values
- **Caution: This can happen easily when using dynamic partitioning!**
  - For example, partitioning customers by first name could produce thousands of partitions



## Partitioning in Hive (1)

- In older versions of Hive, dynamic partitioning is not enabled by default
  - Enable it by setting these two properties

```
SET hive.exec.dynamic.partition=true;  
SET hive.exec.dynamic.partition.mode=nonstrict;
```

- Note: Hive variables set in Beeline are for the current session only
  - Your system administrator can configure settings permanently



## Partitioning in Hive (2)

---

- Caution: if the partition column has many unique values, many partitions will be created
- Three Hive configuration properties exist to limit this
  - **hive.exec.max.dynamic.partitions.pernode**
    - Maximum number of dynamic partitions that can be created by any given node involved in a query
    - Default 100
  - **hive.exec.max.dynamic.partitions**
    - Total number of dynamic partitions that can be created by one HiveQL statement
    - Default 1000
  - **hive.exec.max.created.files**
    - Maximum total files (on all nodes) created by a query
    - Default 100000

# Chapter Topics

## Data File Partitioning

## Importing and Modeling Structured Data

- Partitioning Overview
- Partitioning in Impala and Hive
- **Conclusion**
- Hands-On Exercise: Partition Data in Impala or Hive

## Essential Points

---

- Partitioning splits table storage by column values for improved query performance
- Partitions are HDFS directories
  - Names follow the format *column=value*
- Partitions can be defined and loaded dynamically or statically
- Only partition on columns with a reasonable number of possible values

## Bibliography

---

**The following offer more information on topics discussed in this chapter**

- **Impala documentation on partitioning**
  - <http://tiny.cloudera.com/impalapart>
- **Improving Query Performance Using Partitioning in Apache Hive (Cloudera Engineering Blog)**
  - <http://tiny.cloudera.com/partblog>

## Chapter Topics

### Data File Partitioning

### Importing and Modeling Structured Data

- Partitioning Overview
- Partitioning in Impala and Hive
- Conclusion
- **Hands-On Exercise: Partition Data in Impala or Hive**

## Hands-On Exercise: Partition Data in Impala or Hive

---

- **In this exercise you will**
  - Create a table for accounts that is partitioned by area code
- **Please refer to the Hands-On Exercise Manual for instructions**